

WIRSCHING’S WEAK COVERING CONJECTURE: AN EXTENDED COMPUTATION, A CONDITIONAL BOUND, AND THE STRUCTURE OF WHAT RESISTS IT

RENATO AUGUSTO TAVARES

Universidade Federal de Goiás

rat@discente.ufg.br

ORCID: <https://orcid.org/0009-0002-0196-3311>

ABSTRACT. Wirsching’s 1998 Weak Covering Conjecture asks for a sub-exponential $K(l)$ such that a large enough set $R_{j-1,j} \subset \mathbb{Z}_3^*$, built from strictly decreasing exponent tuples, is forced to cover every unit residue modulo 3^l . We extend the exact computation of $j^*(l)$, the smallest such covering budget, from $l = 20$ to $l = 23$, and compare four growth models for $e(l) := j^*(l) - l \log_4 3$ against the extended table: a constant model is unconditionally impossible (we prove $e(l) \geq \frac{1}{4} \log_2 l - O(1)$) and disfavored at every statistic checked; the growing alternatives are not distinguished from one another by these criteria. We give the best bound known to us, $j^*(l) \leq \frac{9}{8}l + \frac{33}{2}$, via a mean-payoff-game relaxation of the covering budget whose proof rests on one computationally verified, not generally proven, correspondence. We then examine the natural Fourier approach to a tighter bound: the natural full-spectrum positivity criterion is provably unreachable at every budget, since no element of $R_{j,k}$ is divisible by 3; granting the criterion regardless, a uniform square-root-cancellation argument still could not beat a threshold worse than the mean-payoff bound. Separately, a direct computation at one accessible level finds the bulk of the primitive-frequency mass sitting outside the sparse exceptional set a fixed magnitude threshold defines. The residues that actually fail to be covered sit in cells of local hit-density well below the global mean; a phase-randomization experiment at a different, more tractable pair of parameters points the same way, toward phase structure that magnitude alone does not capture. We prove a theorem on the uncovered residues themselves, for budgets $j \geq l$: the set left uncovered at a budget is contained in both twice and four times the set left uncovered at the previous budget, giving an explicit recursive upper bound on $j^*(l)$ that we prove exact at most budgets, given one further combinatorial property verified for $l = 3, \dots, 13$ against a table running to $l = 23$; at the one remaining budget, this holds at four small levels checked directly and is otherwise only empirically exact. We falsify the natural cost-1 form of a local repair rule this theorem might suggest, with explicit witness-level exhaustive counts, prove that every last-uncovered residue is $\equiv 1 \pmod{3}$ whenever $j^*(l) \geq l + 1$ (every computed level but one), prove a two-class containment law modulo 9 on the same residues, and report further exact identities, some verified but not proven, governing the last residues to be covered. Two of the questions this work leaves open are named and stated as exact conditions, each narrower than the conjecture itself; others are noted where they arise. The Weak Covering Conjecture remains open.

Keywords: Collatz map, 3-adic dynamics, covering conjecture, mean-payoff games, exponential sums, Littlewood–Offord theory.

2020 Mathematics Subject Classification: 11B50, 11K70, 91A50, 11L07, 37A44.

1. INTRODUCTION

Wirsching’s 1998 monograph [1] studies the density of predecessor sets under the $3n + 1$ map through an averaging construction on the 3-adic integers. A limiting Markov chain built from that construction has an invariant density, and Wirsching’s own route to positive predecessor density, developed further in a later paper of his own that isolates it as a self-contained target and states, as open conjectures, the gaps still unbridged between that construction and a proof [2], passes through a covering question about a family of finite subsets of $\mathbb{Z}_3^*$, the 3-adic units. For integers $j, k \geq 0$,

write

$$(1) \quad R_{j,k} := \left\{ \sum_{i=0}^j 2^{a_i} 3^i \mid j+k \geq a_0 > a_1 > \cdots > a_j \geq 0 \right\} \subset \mathbb{Z}_3^*.$$

Each element of $R_{j,k}$ is built from a strictly decreasing sequence of $j+1$ “digit positions” chosen from $\{0, \dots, j+k\}$: the i -th chosen position, counting from the top, contributes 2^{a_i} scaled by 3^i . Wirsching’s Weak Covering Conjecture (Conjecture 3.9, Chapter V, [1]) asks for a sub-exponential $K(l)$, meaning $\lim_{l \rightarrow \infty} K(l)e^{-\gamma l} = 0$ for every constant $\gamma > 0$ (Wirsching’s own definition, same chapter), such that $|R_{j-1,j}| \geq K(l) \cdot 3^l$ already forces $R_{j-1,j}$ to cover every unit residue modulo 3^l ; the cardinality on the left is that of the tuple family itself, before reduction, which by Proposition 2’s bijection equals $\binom{2j}{j}$; the (a priori much smaller) size of its image modulo 3^l is a different quantity.

Writing $j^*(l)$ for the smallest j such that $R_{j-1,j}$ covers $(\mathbb{Z}/3^l\mathbb{Z})^*$ exactly (a quantity that exists for every l , unconditionally; Proposition 16 below), the conjecture’s forcing form and the growth rate of $j^*(l)$ coincide once coverage, once achieved, persists at every larger budget: this holds for $j \geq l$ (Theorem 10, Section 6), which covers every level in Table 1, but is not established below l , the regime the conjectured rate would eventually enter. With that proviso, the conjecture concerns the growth rate of

$$(2) \quad e(l) := j^*(l) - l \log_4 3,$$

the excess of the true covering budget over the leading-order rate a cardinality argument alone predicts ($|R_{j-1,j}| = \binom{2j}{j} \sim 4^j / \sqrt{\pi j}$, so matching $2 \cdot 3^{l-1}$ residues needs $j \sim l \log_4 3$ at leading order).

At the threshold, $K(l) \asymp |R_{j^*(l)-1,j^*(l)}|/3^l \asymp 4^{e(l)} / \sqrt{j^*(l)}$; the $\sqrt{j^*(l)}$ correction is polynomial in l at every level actually computed, where $j^*(l) = O(l)$ (Table 1), though no unconditional polynomial bound on $j^*(l)$ is proved here (Proposition 16’s unconditional bound is exponential; Corollary 6’s linear bound is conditional). Since a polynomial correction never changes whether a sequence is sub-exponential, this identification, like the rest of this paragraph, carries the persistence proviso above: with that understood, the conjecture’s truth would force $e(l)$ to grow strictly slower than linearly. The converse, that $e(l) = o(l)$ would in turn establish the conjecture, needs coverage to persist at every budget past $j^*(l)$, which Theorem 10 only gives for $j \geq l$; below l , the regime a truly sub-linear $e(l)$ would eventually reach, the converse is open.

An earlier, unpublished manuscript of the author computed $j^*(l)$ exactly for $l = 1, \dots, 20$ by direct enumeration and reported the resulting table without a growth-rate analysis. This paper does three things. First, it extends that table to $l = 23$ and gives a proper statistical comparison of candidate growth models for $e(l)$ against the full range (Sections 2 and 3). Second, it gives the best upper bound on $j^*(l)$ we have found, via a reformulation of the covering budget as a mean-payoff game (Section 4). Third, and at greatest length, it reports a sustained attempt to close the gap to the conjectured rate directly, first by the naive uniform-cancellation version of the natural Fourier approach, and, once that version is shown to fail for a checkable reason, by a direct study of the residues that actually resist coverage (Sections 5–7). That second attempt does not settle the conjecture. It produces a theorem on those residues, a falsification of the cost-1 form of a repair rule the theorem might suggest, and two named open questions, each narrower than the conjecture itself and worth a dedicated attack in its own right.

2. THE COMPUTATION

Table 1 gives $j^*(l)$ and $e(l)$ for $l = 1, \dots, 23$. The values for $l \leq 20$ reproduce the earlier manuscript’s table exactly; they were independently recomputed here by a from-scratch Rust reimplement of the covering search (a bitset rotation dynamic program over residues modulo 3^l), validated against the original Python implementation at every $l \leq 20$ before being trusted for the extension. Each reported $j^*(l)$ is the smallest budget at which the DP’s own coverage bitset is full,

so the search certifies both coverage at $j^*(l)$ and its absence at $j^*(l) - 1$ internally, by construction, at every level it runs; the $l \leq 20$ agreement against an independently written implementation is this paper's evidence that the search itself is correct; it is carried forward by inference for $l = 21, 22, 23$, without a separate cross-check at those levels. The values for $l = 21, 22, 23$ are new.

The search above only tries budgets $j \geq l$ at level l , so on its own it certifies $j^*(l)$ only against smaller candidates within that range.

Lemma 1 (No smaller budget covers). *For $l = 2, \dots, 23$, no $j < l$ has $R_{j-1,j}$ covering $(\mathbb{Z}/3^l\mathbb{Z})^*$, so $j^*(l) \geq l$ holds outright, not only relative to the range the search above tries.*

Proof. If $R_{j-1,j}$ covers $(\mathbb{Z}/3^L\mathbb{Z})^*$ for some $L \geq j$, it covers $(\mathbb{Z}/3^j\mathbb{Z})^*$ too: reduction modulo 3^j factors through reduction modulo 3^L , and the ring homomorphism $\mathbb{Z}/3^L\mathbb{Z} \rightarrow \mathbb{Z}/3^j\mathbb{Z}$ restricts to a surjection of unit groups (every unit modulo 3^j lifts to a unit modulo 3^L), so the image of a full set of units modulo 3^L is a full set of units modulo 3^j .

$R_{0,1} = \{1, 2\}$ has 2 elements, short of $|(\mathbb{Z}/3^l\mathbb{Z})^*| = 2 \cdot 3^{l-1} \geq 6$ for $l \geq 2$, so $j = 1$ never covers at $l \geq 2$: $j^*(l) \geq 2$ for $l \geq 2$.

Fix l with $2 \leq l \leq 23$ and suppose, for contradiction, some j_0 with $2 \leq j_0 < l$ has R_{j_0-1,j_0} covering $(\mathbb{Z}/3^l\mathbb{Z})^*$. The reduction above, with $L = l$, gives R_{j_0-1,j_0} covering $(\mathbb{Z}/3^{j_0}\mathbb{Z})^*$ too, i.e. $j^*(j_0) \leq j_0$. But $j_0 < l \leq 23$, so j_0 is itself a level the search above already ran, at budgets starting from j_0 itself; Table 1 reports $j^*(j_0) > j_0$ strictly there. Contradiction. So no $j_0 < l$ covers, giving $j^*(l) \geq l$. $\square$

l	$j^*(l)$	$e(l)$	l	$j^*(l)$	$e(l)$	l	$j^*(l)$	$e(l)$
1	1	0.208	9	13	5.868	17	21	7.528
2	4	2.415	10	15	7.075	18	22	7.735
3	6	3.623	11	16	7.283	19	23	7.943
4	7	3.830	12	17	7.490	20	24	8.150
5	9	5.038	13	18	7.698	21	25	8.358
6	10	5.245	14	19	7.905	22	26	8.565
7	11	5.453	15	20	8.113	23	27	8.773
8	12	5.660	16	20	7.320			

TABLE 1. $j^*(l)$ and $e(l) = j^*(l) - l \log_4 3$, $l = 1, \dots, 23$.

The single plateau in the table, $j^*(15) = j^*(16) = 20$, is not an error: $l = 15 \rightarrow 16$ is the only step among the 22 from $l = 1$ to $l = 23$ at which the covering budget does not increase, and it is exactly the event Section 3's plateau-frequency test measures. The integer difference $j^*(l) - l$, a coarser cousin of $e(l)$, is 0, 2, 3, 3, 4, 4, 4, 4, 4, 5, 5, 5, 5, 5, 4, 4, 4, 4, 4, 4, 4, 4 for $l = 1, \dots, 23$: it never grows past 5 over the computed range, and nothing proven here rules out that ceiling persisting; Proposition 2's logarithmic bound on $e(l)$ does not touch it, since a bounded $j^*(l) - l$ makes $e(l) = (j^*(l) - l) + l(1 - \log_4 3)$ grow linearly, comfortably above a logarithmic floor. If the ceiling did persist, $j^*(l) = l + O(1)$ would make $e(l)$ linear, and, by the identity noted after (2), would (subject to the persistence proviso stated there) make $K(l)$ exponential and falsify the conjecture. This flatness is a fact about the range checked, not a proven bound either way.

$l = 21$ fits in physical RAM once a memory-wasteful parallelization step in the original implementation was corrected. $l = 22$ needs approximately 84 GiB of state, above the machine's 62 GiB of physical RAM, so a 500 GiB swap file was added to the machine specifically for this computation; the run took approximately 10,749 s (~ 3 hours) with swap I/O as the dominant cost, against 748 s for $l = 21$ without swap. $l = 23$ required approximately 263 GiB of state, comfortably within the same swap budget, and 375,615 s (~ 104 hours). Three attempts at $l = 24$, made over the course of this project on the same machine, each failed before completion (two to memory-policy kills, one

to an apparent out-of-memory condition with no surviving kernel log confirming it directly); $l = 24$ is not attempted further, and the table stops at $l = 23$.

3. THE GROWTH OF $e(l)$

We compare four candidate functional forms for $e(l)$'s growth against the tail $l = 10, \dots, 23$ ($n = 14$; the same range choice as the earlier manuscript, now with three more points), each fit by ordinary least squares: $e(l) = c$ (constant), $e(l) = a + b \ln l$ (logarithmic), $e(l) = a + b\sqrt{l}$ (square-root), and $e(l) = a + bl$ (slow-linear). Table 2 gives the fit statistics, using the Gaussian-OLS forms $AIC = n \ln(RSS/n) + 2k$, $BIC = n \ln(RSS/n) + k \ln n$ (RSS the residual sum of squares, k the number of regression coefficients, 1 for the constant model and 2 for the other three), and leave-one-out cross-validated RMSE; since only differences ΔAIC , ΔBIC are used below, the convention for a shared additive constant across models does not affect them.

model	k	ΔAIC	ΔBIC	LOOCV RMSE
constant	1	16.091	15.452	0.5200
logarithmic	2	1.512	1.512	0.2991
square-root	2	0.722	0.722	0.2909
slow-linear	2	0.000	0.000	0.2837

TABLE 2. Model comparison on $e(l)$, $l = 10, \dots, 23$. $\Delta AIC/BIC$ relative to the best-fitting model (slow-linear).

A constant model is disfavored sharply and consistently: $\Delta AIC = 16.09$, $\Delta BIC = 15.45$, both growing at every l added since $l = 21$. The three growth models are not distinguished from one another by the conventional $\Delta AIC > 2$ rule of thumb, with slow-linear taking a narrow lead on both AIC and leave-one-out cross-validation; the three candidate regressors ($\log l$, $\sqrt{l}$, l) are themselves highly correlated over this range (pairwise correlation ≥ 0.993), so this ranking mainly reflects how thin the current margin is. By the identity noted after (2), a linear $e(l)$ would make $K(l)$ exponential, falsifying the conjecture outright. Fourteen highly correlated points cannot distinguish that possibility from logarithmic or square-root growth, so this comparison does not settle the question either way; it only says which growth rates remain live.

None of this machinery is actually needed to rule out a constant $e(l)$: since $\log_4 3 < 1$, rounding a constant-plus-linear quantity $l \log_4 3 + c$ to the nearest integer can only change by 0 or 1 from one l to the next, yet $j^*(1) = 1$ and $j^*(2) = 4$, an increment of 3. A constant- $e(l)$ source is incompatible with the table's first two entries alone, with no statistical test required.

A direct combinatorial test tells the same story from a different angle, restricted to where its own null hypothesis is not already dead: the constant-rounding model behind it predicts every increment in $\{0, 1\}$, which the full table already refutes (the increments of 2 and 3 noted above), so a plateau count over the full range tests nothing about that model. Restricted to the $l = 10, \dots, 22$ tail the model comparison above uses (13 increments over the same 14 points, all actually in $\{0, 1\}$, consistent with the model's own prediction), exactly one is zero (the $l = 15 \rightarrow 16$ plateau noted above); the constant-rounding model predicts a plateau at a fixed fraction $1 - \log_4 3 \approx 0.2075$ of increments (the fraction of the real line, in the limit, that rounds down rather than up under a linear-plus-constant quantity with irrational slope $\log_4 3$); testing the observed count of 1 against $\text{Bin}(13, 1 - \log_4 3)$, one-sided, gives $p \approx 0.214$. $e(l)$ is an exact, deterministic sequence, so this carries no sampling interpretation; it is best read descriptively, as "how well does a constant- $e(l)$ source fit the observed tail increment pattern." Under the constant-rounding model, $e(l+1) - e(l)$ can only be $1 - \log_4 3 \approx 0.2075$ or $-\log_4 3 \approx -0.7925$; the plateau count above and Section 3's OLS comparison are both functions of the same 14-point increment sequence, so they are not independent checks,

though the plateau test bears specifically on the constant-rounding model, a different object from the free-intercept constant regression $e(l) = c$ that Table 2 fits.

A constant $e(l)$ is disfavored descriptively by every measure above. It is also, unconditionally, impossible.

Proposition 2 (Unboundedness of $e(l)$). $e(l) \geq \frac{1}{4} \log_2 l - O(1)$.

Proof. Distinct exponent tuples give distinct sums: since the a_i are distinct, a_{j-1} is the unique minimal one, so $V = \sum_i 2^{a_i} 3^i$ has 2-adic valuation exactly a_{j-1} (every other term is divisible by a strictly higher power of 2), recovering a_{j-1} from V directly. The full term $2^{a_{j-1}} 3^{j-1}$ is then subtracted off, and the process iterates on the remaining $j-1$ terms. So $R_{j-1,j}$, in bijection with j -subsets of the $2j$ -element set $\{0, \dots, 2j-1\}$, has exact cardinality $\binom{2j}{j}$ (this cardinality is Wirsching's own Corollary 1.8, Chapter V, [1]; the injectivity argument above is this paper's). Covering $(\mathbb{Z}/3^l\mathbb{Z})^*$, of size $2 \cdot 3^{l-1}$, requires the domain to be at least as large as the codomain: $\binom{2j}{j} \geq 2 \cdot 3^{l-1}$ is necessary at $j = j^*(l)$ in particular, since $R_{j^*(l)-1, j^*(l)}$ covers by definition of $j^*(l)$. The explicit bound $\binom{2n}{n} \leq 4^n / \sqrt{3n+1}$, valid for every $n \geq 1$, turns this into $4^j \geq 2 \cdot 3^{l-1} \sqrt{3j+1}$, hence

$$j \geq l \log_4 3 - \log_4 3 + \log_4 2 + \frac{1}{2} \log_4 (3j+1).$$

Dropping the last, non-negative term already gives $j \geq l \log_4 3 - O(1)$, so $3j+1 = \Omega(l)$, hence $\log_4 (3j+1) \geq \log_4 l - O(1) = \frac{1}{2} \log_2 l - O(1)$. Feeding this back into the displayed inequality gives the claimed bound. $\square$

4. A MEAN-PAYOFF-GAME UPPER BOUND

The full-precision game. Fix $l \geq 1$. A *full-precision play* from a unit z_0 modulo 3^l is a sequence of l moves: at step $i = 0, \dots, l-1$, from a unit z_i modulo 3^{l-i} , a move of cost $d_i \geq 0$ is *legal* when $2^{d_i} z_i \equiv 1 \pmod{3}$, and produces $z_{i+1} := T_{d_i}(z_i) := (2^{d_i+1} z_i - 2)/3$, which is then required to be a unit modulo 3^{l-i-1} (automatic at $i = l-1$, where $3^{l-i-1} = 1$). Write $J := \sum_{i=0}^{l-1} d_i$ for the total cost.

Lemma 3 (Grounding). *For a legal full-precision play from z_0 with costs $d_0, \dots, d_{l-1}$ and total cost J , define $a_{l-1} := 0$ and $a_{i-1} := a_i + d_i + 1$ for $i = l-1, \dots, 1$ (so $a_0 = J - d_0 + l - 1$). Then*

$$2^{J+l-1} z_0 \equiv \sum_{i=0}^{l-1} 2^{a_i} 3^i \pmod{3^l}.$$

The right side is, by construction, an element of $U(l, a_0)$ (Section 7) with bottom exponent $a_{l-1} = 0$.

Proof. By induction on the play length n ($n = 1, \dots, l$), stated for a play of length n starting at a unit z modulo 3^n (matching the lemma's notation at $n = l$, $z = z_0$). Write $J_n := \sum_{i=0}^{n-1} d_i$ and $a_{n-1} := 0$, $a_{i-1} := a_i + d_i + 1$ for $i = n-1, \dots, 1$, as above.

Base case $n = 1$: the single move has $2^{d_0} z \equiv 1 \pmod{3}$, which is exactly $2^{a_0+d_0} z \equiv 2^{a_0} 3^0 \pmod{3}$ since $a_0 = 0$ when $n = 1$.

Inductive step: suppose the claim holds at length $n-1$ for the sub-play with costs $d_1, \dots, d_{n-1}$ starting at $z' := T_{d_0}(z)$, a unit modulo 3^{n-1} ; writing $a'_i := a_{i+1}$ for this sub-play, $a'_{n-2} = a_{n-1} = 0$ matches its base case, and its recursion $a'_{i-1} = a'_i + d_{i+1} + 1$ matches $a_i = a_{i+1} + d_{i+1} + 1$ (the original recursion at index $i+1$), so the induction hypothesis reads

$$2^{a_1+d_1} z' \equiv \sum_{i=0}^{n-2} 2^{a_{i+1}} 3^i \pmod{3^{n-1}}.$$

By definition of T_{d_0} , $2^{d_0+1} z \equiv 3z' + 2 \pmod{3^n}$. Multiplying by the unit $2^{a_1+d_1}$ gives

$$2^{a_1+d_1+d_0+1} z \equiv 3 \cdot 2^{a_1+d_1} z' + 2^{a_1+d_1+1} \pmod{3^n}.$$

Since $a_0 = a_1 + d_1 + 1$ (the recursion at $i = 1$), the left side is $2^{a_0+d_0}z$ and the last term is 2^{a_0} . Multiplying the induction hypothesis by 3 turns it, modulo 3^n , into $3 \cdot 2^{a_1+d_1}z' \equiv \sum_{i=1}^{n-1} 2^{a_i}3^i$, so

$$2^{a_0+d_0}z \equiv 2^{a_0} + \sum_{i=1}^{n-1} 2^{a_i}3^i = \sum_{i=0}^{n-1} 2^{a_i}3^i \pmod{3^n},$$

the claim at length n , using $a_0 + d_0 = J_n + n - 1$ (telescoping the recursion). Verified directly against every legal play for $l = 1, \dots, 4$ as a check, with no exception, before being trusted. $\square$

Empirical Result 4 (The game's worst case is $j^*(l)$).

$$j^*(l) = \max_{z_0} \min\{J : z_0 \text{ has a legal full-precision play of total cost } J\},$$

the maximum over units z_0 modulo 3^l . Verified against the known exact table of $j^*(l)$ for $l = 1, \dots, 12$, with no discrepancy; that a legal full-precision play exists at all from every unit z_0 (the minimum inside the display is over a nonempty set) is part of what this verification confirms empirically. A general proof would additionally need the successor to stay a unit at the next modulus at each step, a condition the mod-3 legality check alone does not decide. Lemma 3 does not by itself establish this: the twist 2^{J+l-1} in the lemma depends on z_0 through its own optimal J , so it is not a single, z_0 -independent relabeling of the unit group, and we do not have a general proof of the displayed identity.

The window- k relaxation. A move of cost $d \geq 0$ at z is *legal* when $2^d z \equiv 1 \pmod{3}$, as above, and *safe* at window k when, for $m := 3^k$, every one of the three residues $z, z + m, z + 2m$ modulo $3m$ maps under T_d back into a unit modulo m . A minimizing player, seeing only $z \bmod 3^k$, chooses a legal, safe d ; an adversary then reveals the hidden next 3-adic digit $\epsilon \in \{0, 1, 2\}$, forcing the true state to $T_d(z + 3^k \epsilon) \bmod 3^k$. Both requirements matter: legality alone can send a unit to 0 (take $z = 1, d = 0$: $T_0(1) = 0$), and the safety filter is exactly what keeps the walk inside $\mathbb{Z}_3^*$ regardless of which digit the adversary supplies. Legality depends on d only through $d \bmod 2$ (since 2 has order 2 modulo 3), and 2 is a primitive root modulo 3^{k+1} for every $k \geq 0$ (order $2 \cdot 3^k$, by an elementary lifting-the-exponent induction), so $2^{d+1} \bmod 3^{k+1}$, and hence $T_d(z) \bmod 3^k$ and the safety check itself, depend on d only through $d \bmod (2 \cdot 3^k)$: successor states repeat with this period as d grows, and every representative beyond the smallest in its residue class has strictly larger cost for the same effect, so is never used by an optimal policy. The action set at each state is therefore, without loss, the subset of $\{0, \dots, 2 \cdot 3^k - 1\}$ that is legal and safe there, which is finite. That this subset is nonempty at every unit state has been verified computationally at every $k = 3, \dots, 14$ used below; a general proof is left open here.

This restricted-visibility game is, by construction, a two-player *mean-payoff game* on the $2 \cdot 3^{k-1}$ unit states modulo 3^k , with the now finite legal, safe moves as the minimizer's edges, the adversary's three choices of ϵ as the maximizer's edges, and payoff the long-run average of the move costs d . The classical theory of such games, initiated by Ehrenfeucht and Mycielski [5], guarantees a value and an optimal positional policy in general; Theorem 5 below needs none of that generality, only the twelve specific policies whose resulting (ρ_k, C_k) are reported in Table 3, each independently checked as described in Section 9.

A window- k policy, run against the actual digit sequence, gives a legal full-precision play at any z modulo 3^l : at step i , the true remaining state is z_i modulo 3^{l-i} , so the policy sees the true $z_i \bmod 3^k$ directly whenever $l-i \geq k$, i.e. for every step but the last $k-1$. Once $l-i < k$, whether because $l < k$ globally or because a play with $l \geq k$ has reached its own final $k-1$ steps, fix any lift of the true state to a residue modulo 3^k and run the policy on the lift, which is exactly the situation safety was built for, with the play's own remaining digits standing in for the hidden ones, so the chosen move stays legal against the true, lower-precision state at every step. Writing s_i for whichever window- k state σ_k sees this way at step i (the true $z_i \bmod 3^k$, or a fixed lift of it), s_{i+1} is

by construction the successor $s'_i(\epsilon)$ under σ_k 's chosen move and the digit ϵ that the real play actually supplies at that step, so $s_0, \dots, s_l$ is exactly a play of the form the proof below sums over, with move costs equal to the real play's (this regime is part of what the policy-against-the-covering-DP check at $k = 4, 7$ against l up to 14 already exercised, with zero illegal or unsafe moves; every play there passes through this same final- $k - 1$ -steps regime regardless of how large l is).

Theorem 5 (Window relaxation bound, conditional on Empirical Result 4). *For each $k = 3, \dots, 14$, the certified policy underlying Table 3 (a legal, safe move at every unit state modulo 3^k , average cost ρ_k) gives an explicit rational C_k with $j^*(l) \leq \rho_k \cdot l + C_k$ for all l .*

Proof. Fix k in range. The policy σ_k , the value ρ_k , and a potential $h_k : (\mathbb{Z}/3^k\mathbb{Z})^* \rightarrow \mathbb{Q}$ are computed by a nested Howard strategy-improvement solver and checked, state by state and digit by digit over the finite action set established above, against

$$d_{\sigma_k}(s) + h_k(s'(\epsilon)) \leq \rho_k + h_k(s)$$

at every state s and every adversary digit ϵ , where $s'(\epsilon)$ is the successor under σ_k 's chosen move. Summing the inequality along any l -step play $s_0, \dots, s_l$ telescopes to

$$\sum_{i=0}^{l-1} d_{\sigma_k}(s_i) \leq l\rho_k + h_k(s_0) - h_k(s_l) \leq l\rho_k + (\max h_k - \min h_k) =: l\rho_k + C_k,$$

regardless of which digits the adversary actually supplies, since the potential inequality holds for every ϵ at every step. So σ_k , run as above, gives every unit z modulo 3^l a legal full-precision play of total cost at most $l\rho_k + C_k$; the minimum cost over all legal plays from z is therefore also at most $l\rho_k + C_k$. Taking the worst-case z and applying Empirical Result 4 gives $j^*(l) \leq \rho_k l + C_k$. $\square$

σ_k , h_k and ρ_k are computed as rationals by a nested Howard strategy-improvement solver, which searches actions d up to a fixed cap of 40 (well above the largest d any certified policy below actually uses at any $k = 3, \dots, 14$, which is 11). The potential inequality above is checked, state by state and digit by digit, against the full legal, safe action set at that state, so the resulting upper bound $j^*(l) \leq \rho_k l + C_k$ does not depend on the cap. The solver also extracts a matching adversary lower bound from the same run, but that computation ranges only over the same capped action set: “upper bound equals lower bound” certifies ρ_k exactly as the value of the mean-payoff game restricted to $d \leq 40$, not, without further argument, as the value of the unrestricted game on the full action set $\{0, \dots, 2 \cdot 3^k - 1\}$ described above, since a move with $d > 40$ could in principle let the adversary force a lower true value. Table 3 reports ρ_k , C_k under this understanding: certified rates of an exhibited, verified policy, which is all Theorem 5 needs, not a proven-exact value of the full window- k game. σ_k and ρ_k are cross-validated by an independent value-iteration solver and, at $k = 4$ and $k = 7$, by running the resulting policy against the actual finite covering DP for every unit residue modulo 3^l , l up to 14, with zero illegal or unsafe moves encountered. Table 3 gives the values computed for $k = 3, \dots, 14$.

k	3	4	5	6	7	8	9	10	11	12	13	14
ρ_k	2	$\frac{5}{3}$	$\frac{3}{2}$	$\frac{3}{2}$	$\frac{7}{5}$	$\frac{25}{19}$	$\frac{5}{4}$	$\frac{11}{9}$	$\frac{6}{5}$	$\frac{7}{6}$	$\frac{119}{104}$	$\frac{9}{8}$
C_k	5	$\frac{19}{3}$	$\frac{15}{2}$	9	10	$\frac{207}{19}$	$\frac{47}{4}$	$\frac{115}{9}$	$\frac{69}{5}$	$\frac{44}{3}$	$\frac{1621}{104}$	$\frac{33}{2}$

TABLE 3. The window-relaxation value ρ_k and constant C_k , $k = 3, \dots, 14$.

At $k = 14$, $\rho_{14} = 9/8 = 1.125$ and $C_{14} = 33/2$, giving:

Corollary 6 (Conditional on Empirical Result 4). $j^*(l) \leq \frac{9}{8}l + \frac{33}{2}$.

This is, to our knowledge, the best bound on $j^*(l)$'s growth rate currently available, conditional on Empirical Result 4. ρ_k is observed to be non-increasing over the computed range $k = 3, \dots, 14$ (Table 3); we have not verified this in general, and Corollary 6 uses only the single value $k = 14$. The general technique, representing a covering-type combinatorial problem as a mean-payoff game and computing exact rational relaxation bounds this way, is not new to this application: beyond Ehrenfeucht and Mycielski's foundational theorem [5] itself, a close and recent methodological precedent proves rationality and computability of the covering radius of a primitive sofic shift by essentially the same construction, applied to a different object [6]. We cite both as precedent for the framing; what is new here is its specific application to $j^*(l)$.

Remark 7. Fitting $\rho_k = L + A/k$ by ordinary least squares to $k = 10, \dots, 14$ gives $L \approx 0.876$, still well above the conjectured true rate $\log_4 3 = 0.7925\dots$; the fit is sensitive to the range used (all of $k = 3, \dots, 14$ gives $L \approx 0.904$) and is not evidence of any specific limit. Only the direction $j^*(l) \leq \rho_k l + C_k$, proven above, is established. The reverse inequality, $\inf_k \rho_k \leq \limsup_l j^*(l)/l$, which would be needed to conclude anything about where ρ_k is heading, is not established.

5. THE EXPONENTIAL SUM, AND WHY IT RESISTS A DIRECT ATTACK

Corollary 6's conditional bound, $1.125l$, is well above the conjectured rate $\log_4(3)l \approx 0.7925l$. The natural route to closing that gap is Fourier-analytic. Here $e(x) := e^{2\pi i x}$ denotes the standard additive character, unrelated to the sequence $e(l)$ of (2). By orthogonality on $\mathbb{Z}/3^l\mathbb{Z}$, the number of tuples landing in residue class z is

$$(3) \quad N(z) = \frac{1}{3^l} \sum_{t=0}^{3^l-1} S(t) e\left(-\frac{tz}{3^l}\right), \quad S(t) := \sum_{j+k \geq a_0 > \dots > a_j \geq 0} e\left(\frac{t}{3^l} \sum_{i=0}^j 2^{a_i} 3^i\right),$$

and $S(0) = |R_{j,k}| =: T$. The natural route to positivity is the triangle inequality: $N(z) > 0$ for every z as soon as $\sum_{t \neq 0} |S(t)| < T$. This route is closed before it opens. Proposition 8 below shows the sum on the left is never below T , for any l, j, k : the criterion cannot be reached at all, for any budget. The remainder of this section studies what is left once the full-spectrum route is abandoned: a primitive-frequency version of the same idea (§5.3), and why it too resists the reasons below.

5.1. A hard ceiling on uniform cancellation.

Proposition 8 (The full-spectrum criterion is unreachable). *For every $l \geq 1$ and every $R_{j,k}$ with $T := |R_{j,k}| > 0$, $\sum_{t \neq 0} |S(t)| \geq T$; in particular $|S(3^{l-1})| \geq T/2$. Consequently the criterion $\sum_{t \neq 0} |S(t)| < T$ is never satisfied, for any l, j, k , and neither is the uniform bound $|S(t)| = o(T/3^l)$ at every nonzero t , since that would force $\sum_{t \neq 0} |S(t)| = o(T)$.*

Proof. Every element of $R_{j,k}$ reduces to $2^{a_0} \bmod 3$ (all other terms carry a positive power of 3), so it is $\equiv 1$ or $2 \pmod{3}$: no element is divisible by 3. Hence $N(z) = 0$ for every z divisible by 3, in particular $N(0) = 0$, so $\sum_{t=0}^{3^l-1} S(t) = 3^l N(0) = 0$ and $\sum_{t \neq 0} S(t) = -S(0) = -T$; the triangle inequality gives $\sum_{t \neq 0} |S(t)| \geq |\sum_{t \neq 0} S(t)| = T$.

The same collapse localizes: summing the inversion formula over the 3^{l-1} residues divisible by 3 leaves only the three frequencies $t \in \{0, 3^{l-1}, 2 \cdot 3^{l-1}\}$ in the inner sum (every other character is nonconstant on that coset of the index-3 subgroup and averages to 0), giving $S(3^{l-1}) + S(2 \cdot 3^{l-1}) = -T$ exactly, at every l, j, k . Since $S(3^l - t) = \overline{S(t)}$, the two terms are conjugate, so $\operatorname{Re} S(3^{l-1}) = -T/2$ exactly and $|S(3^{l-1})| \geq T/2$: a single conjugate pair already forces the bound, before the rest of the spectrum is added in. $\square$

Proposition 8 closes the naive route outright, at every j : there is no threshold past which it starts to work. Even granting the unreachable premise for comparison, specializing $S(t)$, T to the family $R_{j-1,j}$ this paper's covering question concerns $(T = \binom{2j}{j})$ and asking what j a uniform bound $|S(t)| \leq C\sqrt{T}$ (square-root cancellation), applied via the triangle inequality across all $3^l - 1$ nonzero frequencies, would force gives $(3^l - 1)C\sqrt{T} < T$, so $T > ((3^l - 1)C)^2$, i.e. $4^j/\sqrt{3j+1} \geq T > ((3^l - 1)C)^2$ (the explicit bound from Proposition 2's proof), hence $j \geq 2\log_4(3)l - O(1) \approx 1.585l - O(1)$: already *worse* than Corollary 6's $1.125l$, on top of resting on a premise Proposition 8 rules out at every j . Both facts together rule out the naive version of the Fourier approach, and direct attention to whether a finer, frequency-dependent treatment (a major/minor-arc split, in the language of the circle method) can do better.

5.2. A computation: the exceptional mass does not stay small. Numerically, $|S(t)|$ is far from uniform: the frequencies of largest magnitude sit close to dyadic rationals of small denominator, $t/3^l \approx p/2^r$ for small r , matching exactly the conjugate pair Proposition 8 exhibits at $t = 3^{l-1}, 2 \cdot 3^{l-1}$ ($r = 0$). The natural repair is to restrict to primitive frequencies, excising this pair along with the rest of the non-primitive spectrum, and control the remainder by an averaged bound; whether that remainder stays sparse enough at the scale that matters is a question this section answers computationally, for the family $R_{j-1,j}$ at $l = 18$, $j = 16$ (the first-order cardinality threshold, $T := \binom{32}{16}$, six budgets below the actual covering threshold $j^*(18) = 22$: coverage necessarily fails here already, and this level sits below the $j \geq l$ range Sections 6–7 otherwise work in), with a negative result. Write $\|S\|_1 := \sum_{\gcd(t,3)=1} |S(t)|/T$ for the normalized L^1 mass over frequencies coprime to 3. This restriction excludes $t = 3^{l-1}$, whose magnitude $|S(3^{l-1})| \sim T/\sqrt{3}$ is fixed by Proposition 9 and, empirically, is the largest among frequencies checked at accessible l , along with its conjugate $t = 2 \cdot 3^{l-1}$ (since $S(3^l - t) = \overline{S(t)}$, that frequency matches it exactly and is excluded for the same reason). The restriction to $\gcd(t, 3) = 1$ also excludes every nonprimitive frequency, roughly a third of the spectrum, entirely, so $\|S\|_1$ below measures only the primitive-frequency mass. Here $\|S\|_1 = 5226.01$. A fixed threshold makes the visible spikes look sparse, but of the 8,014 primitive frequencies with $|S(t)|/T$ exceeding 0.001, their combined contribution to $\|S\|_1$ is under 132, so over 97% of the primitive-frequency mass at this single computed level lives in an exponentially numerous population of near-average-magnitude coefficients. This is evidence against a sparse-exceptional-set repair among primitive frequencies at this accessible level; it says nothing about nonprimitive frequencies, about every l , and, since $j = 16$ sits six budgets below $j^*(18)$, still less about the mass distribution near the actual covering threshold.

The same clustering near dyadic rationals connects the problem to a 2011 discussion of Terence Tao's, applying Littlewood–Offord and Riesz-product methods to an object built from the same $2^a 3^b$ -type exponential sums, in the closely related setting of a hypothetical non-trivial Collatz cycle [7]; his construction is structurally close to a boundary slice of $R_{j,k}$ with the bottom exponent pinned at 0. His own conclusion, in his own words, is that “once one uses the transcendence theory bound, ... Littlewood–Offord techniques are no longer of much use,” and that he knows “of no way to get a nontrivial separation property between powers of 2 and powers of 3 other than via transcendence theory methods.”

5.3. A direct test of the residues that actually fail. The most direct test examines the residues that actually fail to be covered, one budget step before covering succeeds, i.e. $H(l, j^*(l) - 1)$ (Section 6). For $l = 10, \dots, 15$ this set has size 2, 1, 3, 1, 3, 1 respectively, tiny against a global occupancy mean $\lambda = |R_{j^*(l)-2, j^*(l)-1}|/(2 \cdot 3^{l-1})$ that is itself in the thousands over the same levels (from 1019 at $l = 10$ to 3695 at $l = 15$). Write the *depth- c cell* of a residue z for its class modulo 3^c , and its *depth- c local intensity* $\lambda_c(z)$ for the total hit count within that class divided by 3^{l-c} , the class's size (defined for $c < l$; at $c = l$ the class is $\{z\}$ itself, and $\lambda_l(z) = 0$ identically for any holdout, a degenerate case excluded below). A *non-homogeneous Poisson model* fit to that intensity

assigns z hole probability $e^{-\lambda_c(z)}$. Computed directly from the exact histogram for $l = 10, \dots, 15$, every one of the eleven holdouts across these six levels has λ_c well below the global mean at every depth $c = 8, 9, 10$ with $c < l$ (at $l = 10$, only $c = 8, 9$; λ_c from 2.0 to 108.4, hole probability from e^{-2} to e^{-108}). Assigning probability e^{-108} to an event that occurred falsifies the non-homogeneous Poisson model outright; what survives is the weaker, descriptive claim underneath it: holdouts sit systematically in low-intensity cells, with how far below the mean varying widely by depth and level. A separate diagnostic at a smaller, more tractable pair, $(l, m) = (14, 16)$, for the family $R_{m-1, m}$ (mean hit count 188.5, a budget below the covering threshold $j^*(14) = 19$), holds every $|S(t)|$ fixed and randomizes the phases of conjugate-symmetric pairs. Writing $N_l(z)$ for the exact hit count of residue z and N_{l-1} for the same tuple family's count one level down, restricting to primitive frequencies $3 \nmid t$ gives exactly $\sum_{3 \nmid t} S(t)e(-tz/3^l) = 3^{l-1}(3N_l(z) - N_{l-1}(z \bmod 3^{l-1}))$, a measure of imbalance among the three lifts of each parent residue. Over unit z , the actual ratio of this quantity's maximum absolute value to its root-mean-square is 15.79, while eight independent phase scrambles (fixed seed) give ratios only in $[5.11, 5.24]$, collapsing to what a generic-phase model predicts once the phases are scrambled away: the primitive lift imbalance in the actual data is far more extreme than a generic-phase null would produce, evidence that the true phases are structured rather than random, though this diagnostic is silent on what mechanism produces that structure or on whether it is what decides coverage. At the levels checked, both the raw holdout rarity and the phase-scramble gap point the same way, toward phase structure that magnitude-only methods (uniform bounds, averages, low moments of $S(t)$) do not see.

Proposition 9. $|S(3^{l-1})|/\binom{2m}{m} \rightarrow 1/\sqrt{3}$ as $m \rightarrow \infty$, for the family $R_{m-1, m}$.

Proof sketch. $\sum_i 2^{a_i} 3^i \equiv 2^{a_0} \pmod{3}$, so $S(3^{l-1})/\binom{2m}{m}$ depends only on the limiting distribution of the top exponent a_0 's parity as $m \rightarrow \infty$, which converges to the mixture $\frac{1}{3}\omega + \frac{2}{3}\omega^2$ of the two nontrivial cube roots of unity $\omega = e^{2\pi i/3}$ (the limiting gap $2m - 1 - a_0$ of the maximum of a random m -subset of $\{0, \dots, 2m - 1\}$ is geometric with parameter $1/2$, giving $\Pr[a_0 \text{ even}] \rightarrow 1/3$). This mixture equals $-\frac{1}{2} - \frac{i\sqrt{3}}{6}$, of modulus $1/\sqrt{3}$. $\square$

$t/3^l = 1/3$ exactly, so l drops out of $S(3^{l-1})$ entirely (only $V \bmod 3$ matters), and the limit above holds at every scale m , including along the covering-relevant diagonal $m \sim l \log_4(3)$. It sharpens Proposition 8's exact bound $|S(3^{l-1})| \geq T/2$ to the precise asymptotic ratio $1/\sqrt{3} \approx 0.577$, and confirms that the mass Proposition 8 forces concentrates, at this one frequency at least, exactly where the numerics of §5.2 place it: at the dyadic rational $t/3^l = 1/3$, $r = 0$.

6. A THEOREM ON THE RESIDUES THAT RESIST COVERAGE

Write $H(l, j)$ for the set of units modulo 3^l not covered by $R_{j-1, j}$. By definition of $j^*(l)$ as the least covering budget, $H(l, j^*(l)) = \emptyset$ and $H(l, j) \neq \emptyset$ for every $j < j^*(l)$; whether this extends to " $j \geq j^*(l)$ exactly when $H(l, j) = \emptyset$ " needs emptiness to persist at every larger budget too, which is not automatic from the definition alone. Every result in this section needs $j \geq l$; $j^*(l) \geq l$ holds outright for $l = 1, \dots, 23$ (equality only at $l = 1$), by Lemma 1 above, but the conjectured rate $\log_4 3 < 1$ would eventually put $j^*(l)$ below l , outside this section's reach. Within this range, persistence does hold: Theorem 10 below gives $H(l, j) = \emptyset \Rightarrow H(l, j+1) \subseteq 2\emptyset \cap 4\emptyset = \emptyset$, so " $H(l, j) = \emptyset$ exactly when $j \geq j^*(l)$ " is established for every $j \geq l$. We make no claim about $j < l$.

Theorem 10 (Holdout doubling inclusion). *For $j \geq l$, $H(l, j+1) \subseteq 2H(l, j) \cap 4H(l, j)$.*

Proof. $R_{j-1, j}$ is the set of values $r = \sum_{i=0}^{j-1} 2^{b_i} 3^i$ over j -subsets $\{b_0 > \dots > b_{j-1}\} \subseteq \{0, \dots, 2j-1\}$, and $R_{j, j+1}$ the same with $(j+1)$ -subsets of $\{0, \dots, 2j+1\}$. Fix $r \in R_{j-1, j}$ with exponent set S , and $s \in \{1, 2\}$. Shifting every exponent up by s gives $S+s \subseteq \{s, \dots, 2j-1+s\} \subseteq \{1, \dots, 2j+1\}$, a j -subset whose induced value is $2^s r$ (a uniform shift of exponents preserves relative rank, so each

term's power of 3 is unchanged). Since $S+s$ contains no 0, adjoining 0 as the new smallest exponent gives a $(j+1)$ -subset of $\{0, \dots, 2j+1\}$: it becomes the new rank- j term, contributing $2^0 \cdot 3^j = 3^j$. So this witness represents $2^s r + 3^j$, and $3^j \equiv 0 \pmod{3^l}$ once $j \geq l$. Hence $2R_{j-1,j} \cup 4R_{j-1,j} \subseteq R_{j,j+1} \pmod{3^l}$ for $j \geq l$. So if $z \in H(l, j+1)$, i.e. z is not represented at budget $j+1$, then z lies in neither $2R_{j-1,j}$ nor $4R_{j-1,j}$: both $z/2$ and $z/4$ are unrepresented at budget j , i.e. $z \in 2H(l, j) \cap 4H(l, j)$. $\square$

Corollary 11 (Chain contraction). *For $j \geq l$: if $x, 2x, \dots, 2^{t-1}x \in H(l, j+1)$ are pairwise distinct (all reduced modulo 3^l , $t \geq 1$), then $2^i x \in H(l, j)$ for every $i = -2, \dots, t-2$ (i.e. $x/4, x/2$ and, when $t \geq 2$, $x, 2x, \dots, 2^{t-2}x$), pairwise distinct, a chain of length $t+1$.*

Proof. Apply Theorem 10 to each of the t elements: $2^i x \in H(l, j+1)$ gives $2^{i-1}x, 2^{i-2}x \in H(l, j)$ for $i = 0, \dots, t-1$. The union of these pairs over all i is exactly $\{2^i x : i = -2, \dots, t-2\}$, $t+1$ elements indexed by $t+1$ consecutive powers of 2. Since 2 has order $2 \cdot 3^{l-1}$ modulo 3^l (an elementary lifting-the-exponent fact, used already in Section 4), any $2 \cdot 3^{l-1}$ consecutive powers of 2 times x exhaust every unit modulo 3^l , and any shorter run is pairwise distinct. The hypothesis that $x, 2x, \dots, 2^{t-1}x$ are pairwise distinct already forces $t \leq 2 \cdot 3^{l-1}$, so it remains to rule out equality. If $t = 2 \cdot 3^{l-1}$, those t elements would be every unit modulo 3^l , including every element of the (nonempty) image of $R_{j,j+1}$; but $H(l, j+1)$, the complement of that image, contains no such element. Hence $t < 2 \cdot 3^{l-1}$, and the $t+1$ output elements, a shorter run of consecutive powers, are pairwise distinct. $\square$

Corollary 12 (Run-length bootstrap). *Writing $\text{maxrun}(H)$ for the greatest t such that H contains t pairwise distinct elements of the form $x, 2x, \dots, 2^{t-1}x$ (all reduced modulo 3^l), with $\text{maxrun}(\emptyset) := 0$, $j^*(l) \leq j + \text{maxrun}(H(l, j))$ for every $j \geq l$.*

Proof. Let $r := \text{maxrun}(H(l, j))$ and suppose, for contradiction, $H(l, j+r) \neq \emptyset$; pick $z \in H(l, j+r)$, a chain of length 1 ($t = 1$) in the sense of Corollary 11. Applying Corollary 11 once at each of the r budgets $j+r-1, j+r-2, \dots, j$ (all $\geq l$, since $j \geq l$), a chain of length t at budget $j'+1$ becomes a chain of length $t+1$ at budget j' each time, so after r applications the length-1 chain at $j+r$ becomes a chain of length $r+1$ in $H(l, j)$, contradicting $\text{maxrun}(H(l, j)) = r$. So $H(l, j+r) = \emptyset$, i.e. $j^*(l) \leq j+r = j + \text{maxrun}(H(l, j))$. $\square$

Empirical Result 13. Corollary 12 is empirically exact, $j^*(l) = j + \text{maxrun}(H(l, j))$, at every (l, j) with $l+1 \leq j \leq j^*(l)$ we were able to compute, $l = 5, \dots, 20$ at every such budget and l up to 22 at the single budget $j = l+1$, including a case with $|H(l, j)| > 3 \times 10^6$ ($l = 3, 4$ were not separately computed here, but Proposition 24 below proves the same equality there). (Past $j^*(l)$, $H(l, j) = \emptyset$ and $\text{maxrun}(\emptyset) = 0$, so the displayed equality no longer holds there; the claim is only about budgets up to and including $j^*(l)$.) The lower endpoint $j = l+1$ has a real justification behind it: equality at $j = l$ itself holds for $l = 1, \dots, 5$ but fails from $l = 6$ on (e.g. $\text{maxrun}(H(6, 6)) = 5$, giving $j + \text{maxrun} = 11 \neq 10 = j^*(6)$; $\text{maxrun}(H(7, 7)) = 5$, giving $12 \neq 11 = j^*(7)$); Corollary 11 and Corollary 12 both need $j \geq l$ to apply at all, but that alone does not explain why $j = l$ specifically starts failing at $l = 6$. The sub-range $l+2 \leq j \leq j^*(l)$ is proven at $l = 3, \dots, 13$: Proposition 24 below, which also proves the boundary budget $j = l+1$ at $l = 3, 4, 5, 6$ specifically. From $l = 7$ on, $j = l+1$ remains empirical only, tied to Proposition 22's own open converse.

A stronger, per-witness converse of Theorem 10, that $x/2, x/4 \in H(l, j)$ already forces $x \in H(l, j+1)$, which would in turn make a maximal doubling chain both sufficient and necessary for survival to a given budget, is false as an unrestricted statement: $1547 \in 2H(7, 8) \cap 4H(7, 8)$, yet $1547 \notin H(7, 9)$ (both of 1547's two budget-9 witnesses span the full exponent range, the one case Theorem 10's argument does not directly control). Section 7 returns to a restricted, weaker form of this converse (corner-redundancy, not the per-witness statement this counterexample refutes) that remains open in general and is proven at specific levels.

A natural strengthening of Theorem 10, that every individual witness admits a bounded-cost local repair, is also false in its most optimistic form.

Call the *Durfee depth* of a witness $\{a_0 > \dots > a_{j-1}\} \in R_{j-1,j}$ at budget j the count $\#\{i : a_i \geq j\}$, the number of chosen exponents at or above the budget itself.

Empirical Result 14 (Failure of the cost-1 local repair rule). At the budget transition $(l, j) = (6, 10) \rightarrow (7, 11)$, call a unit modulo 3^7 a *child* and its reduction modulo 3^6 its *parent*; 100 of the 1458 children (all units modulo 3^7) have minimum achievable Durfee depth exceeding their parent's minimum depth by more than one, up to three; at $(7, 11) \rightarrow (8, 12)$, 332 of 4374 children (all units modulo 3^8) do, likewise up to three. A correspondingly exhaustive check defines the *repair cost* of a child residue r with parent residue p as $\min_{S, S'} |S \Delta S'|$, the smallest symmetric difference of exponent sets over every witness $S \in R_{j-1,j}$ of p and every witness $S' \in R_{j,j+1}$ of r ($|S| = j$, $|S'| = j + 1$, so this distance is always odd, ≥ 1 , and the minimum ranges over the full witness fiber at each budget). This finds repair costs up to 7 at both transitions, against a cost of 1 for the overwhelming majority. This is evidence against a repair cost of exactly 1 at these two transitions; it does not establish that no bounded repair cost exists.

This falsifies, in its originally proposed cost-1 form, a natural conjecture this project had earlier entertained (that coverage propagates via a uniform, cost-1 local repair rule budget by budget); its valid surviving content is Theorem 10 and its corollaries above.

7. FURTHER STRUCTURE OF THE LAST HOLDOUTS

The residues in $H(l, j^*(l) - 1)$, the last to be covered, satisfy several further exact laws. For $W \geq l - 1$, write

$$U(l, W) := \left\{ \sum_{i=0}^{l-1} 2^{\beta_i} 3^i \pmod{3^l} \mid W \geq \beta_0 > \dots > \beta_{l-1} \geq 0 \right\},$$

an l -term analogue of $R_{j-1,j}$ built from a fixed term count and a growing exponent ceiling W .

Lemma 15 (Width scaling identity). *For $j \geq l \geq 1$,*

$$(4) \quad R_{j-1,j} \equiv 2^{j-l} U(l, j+l-1) \pmod{3^l}.$$

Proof. Since $3^i \equiv 0 \pmod{3^l}$ for $i \geq l$, the residue of $r = \sum_{i=0}^{j-1} 2^{b_i} 3^i \in R_{j-1,j}$ modulo 3^l depends only on its top l exponents $b_0 > \dots > b_{l-1}$. A decreasing l -tuple $b_0 > \dots > b_{l-1} \geq 0$ with $b_0 \leq 2j - 1$ extends to a full admissible j -tuple in $\{0, \dots, 2j - 1\}$ (appending $j - l$ further, strictly smaller, distinct exponents) exactly when $b_{l-1} \geq j - l$: append $b_{l-1} - 1, \dots, b_{l-1} - (j - l)$ when this holds, and no extension exists when it does not, since fewer than $j - l$ nonnegative integers lie below b_{l-1} . Writing $c_i := b_i - (j - l)$ turns $b_{l-1} \geq j - l$ into $c_{l-1} \geq 0$ and $b_0 \leq 2j - 1$ into $c_0 \leq 2j - 1 - (j - l) = j + l - 1$, and $\sum_i 2^{b_i} 3^i = 2^{j-l} \sum_i 2^{c_i} 3^i$. So the top- l -exponent tuples realizable by $R_{j-1,j}$ are exactly 2^{j-l} times the decreasing l -tuples counted by $U(l, j+l-1)$, giving the claimed identity modulo 3^l . $\square$

Proposition 16 (Existence). *$j^*(l)$ exists for every $l \geq 1$, with $j^*(l) \leq 2 \cdot 3^{l-1} - 1$.*

Proof. $U(l, \cdot)$ is monotone increasing in W : enlarging the range $\{0, \dots, W\}$ only adds admissible exponent tuples. Shifting every exponent of a tuple realizing $u \in U(l, W)$ up by 1 gives a decreasing l -tuple with top exponent at most $W + 1$ and bottom exponent at least $1 > 0$, still admissible, realizing $2u$; hence $2U(l, W) \subseteq U(l, W + 1)$ for every $W \geq l - 1$. The tuple $\beta_i = l - 1 - i$ ($i = 0, \dots, l - 1$) is admissible at $W = l - 1$ and gives $u_0 := \sum_i 2^{l-1-i} 3^i \equiv 2^{l-1} \pmod{3}$, a unit, so $u_0 \in U(l, l - 1)$. Iterating the inclusion, $2^k u_0 \in U(l, l - 1 + k)$ for every $k \geq 0$. Since 2 has order $2 \cdot 3^{l-1}$ modulo 3^l (already used in Section 4), $\{2^k u_0 : 0 \leq k < 2 \cdot 3^{l-1}\}$ is the coset $u_0 \langle 2 \rangle = (\mathbb{Z}/3^l \mathbb{Z})^*$: every unit modulo 3^l appears as $2^k u_0$ for some k in that range, hence lies in $U(l, l - 1 + k) \subseteq U(l, W^\dagger)$, $W^\dagger := l + 2 \cdot 3^{l-1} - 2$, by monotonicity. So $U(l, W^\dagger) = (\mathbb{Z}/3^l \mathbb{Z})^*$, and Lemma 15 at $j = W^\dagger - l + 1 = 2 \cdot 3^{l-1} - 1 \geq l$ gives $R_{j-1,j} \equiv (\mathbb{Z}/3^l \mathbb{Z})^* \pmod{3^l}$, hence $j^*(l) \leq 2 \cdot 3^{l-1} - 1$. $\square$

This bound is exponential, far above the conjectured $\log_4 3 \cdot l$ rate; it establishes only that $j^*(l)$ is a well-defined finite quantity, which the rest of this section and Section 4's mean-payoff-game argument then bound far more tightly.

By Lemma 15, which itself needs $j \geq l$, the smallest $j \geq l$ with $U(l, j + l - 1) = (\mathbb{Z}/3^l\mathbb{Z})^*$ is the smallest $j \geq l$ at which $R_{j-1,j}$ covers; since $j^*(l) \geq l$ for $l = 1, \dots, 23$, with equality only at $l = 1$ (Lemma 1), this coincides with $j^*(l)$ itself throughout the range this paper works in, though nothing here rules out coverage at some unchecked $j < l$ for a level not yet computed. Writing $W^*(l) := j^*(l) + l - 2$ for the width one step before that,

$$(5) \quad H(l, j) = 2^{j-l} D(l, j + l - 1), \quad D(l, W) := (\mathbb{Z}/3^l\mathbb{Z})^* \setminus U(l, W).$$

Theorem 17 (Last-holdout parity). *If $j^*(l) \geq l+1$, every element of $H(l, j^*(l)-1)$ is $\equiv 1 \pmod{3}$. This covers $l = 2, \dots, 23$ (Table 1); the hypothesis is inapplicable only at $l = 1$ ($j^*(1) = 1$, outside (5)'s range), which says nothing about whether the conclusion itself would hold there, and would become inapplicable again for any l at which the conjectured rate $\log_4 3 < 1$ eventually pushes $j^*(l)$ below $l+1$.*

Proof. $U(l, \cdot)$ is monotone increasing in W , so $D(l, \cdot)$ is monotone decreasing; every $x \in D(l, W^*(l))$ is, by minimality of $j^*(l)$, representable at width $W^*(l) + 1$ but not at $W^*(l)$. Any witness for x at width $W^*(l) + 1$ using only exponents $\leq W^*(l)$ would already witness x at that smaller width, so every such witness has top exponent $\beta_0 = W^*(l) + 1$. Its contribution to the value, at rank 0, is $2^{\beta_0} \cdot 3^0 = 2^{W^*(l)+1}$, so $x \equiv 2^{W^*(l)+1} \equiv (-1)^{W^*(l)+1} \pmod{3}$. By (5), which needs $j^*(l) - 1 \geq l$, elements of $H(l, j^*(l) - 1)$ are $2^{j^*(l)-1-l}x$ for such x , hence

$$2^{j^*(l)-1-l}x \equiv (-1)^{j^*(l)-1-l} \cdot (-1)^{W^*(l)+1} = (-1)^{2j^*(l)-2} = 1 \pmod{3},$$

using $W^*(l) + 1 = j^*(l) + l - 1$ and that $2j^*(l) - 2$ is always even. $\square$

Proposition 18 (Near-extinction: forward containment). *If $j^*(l) \geq l+2$, $H(l, j^*(l)-1) \subseteq 2 \cdot \{x \in H(l, j^*(l)-2) : x \equiv 2 \pmod{3}\}$.*

Proof. Write $J := j^*(l)$. Theorem 10 at $j' = J - 2 \geq l$ gives $H(l, J-1) \subseteq 2H(l, J-2)$, so every $y \in H(l, J-1)$ equals $2x$ for a (unique, since 2 is a unit) $x \in H(l, J-2)$. By Theorem 17, $y \equiv 1 \pmod{3}$; since $y = 2x$, $2x \equiv 1 \pmod{3}$, i.e. $x \equiv 2 \pmod{3}$. So x lies in $\{x \in H(l, J-2) : x \equiv 2 \pmod{3}\}$ and $y = 2x$ lies in twice that set. $\square$

Empirical Result 19 (Near-extinction bijection). $H(l, j^*(l)-1) = 2 \cdot \{x \in H(l, j^*(l)-2) : x \equiv 2 \pmod{3}\}$, exactly, verified at every $l = 2, \dots, 9$ against the exact table with no exception. Proposition 18 proves the $\subseteq$ direction unconditionally (at every level with $j^*(l) \geq l+2$, true throughout this range); only the reverse containment, that every $x \equiv 2 \pmod{3}$ in $H(l, j^*(l)-2)$ actually survives via $2x$, remains empirical.

The same width mechanism that proves Theorem 17 governs the reverse containment one step earlier, through the second-highest exponent β_1 of the extremal witness; unlike the top-exponent argument above, pinning β_1 down needs the one-step width identity of Lemma 23 below. We have not carried that argument through to a full proof of this direction and record it as empirically verified only.

A natural cross-level congruence relating $H(l+1, j^*(l+1)-1)$ to $H(l, j^*(l)-1)$ does not hold. Writing π_l for reduction modulo 3^l , the candidate congruence $x \equiv 4\pi_l(x) \pmod{3^{l+1}}$ is impossible whenever $\pi_l(x)$ is a unit and $l \geq 2$: reducing both sides modulo 3^l forces $\pi_l(x) \equiv 4\pi_l(x) \pmod{3^l}$, i.e. $3\pi_l(x) \equiv 0 \pmod{3^l}$. Checked exhaustively at $l = 2, \dots, 6$, $\pi_l(x)$ never lies in $H(l, j^*(l)-1)$ for $x \in H(l+1, j^*(l+1)-1)$ in the first place, so the pattern has no witnessing instance at these levels either. No cross-level relation between $H(l+1, \cdot)$ and $H(l, \cdot)$ is established by this paper.

Proposition 20 (Mod-9 two-class containment). *For $l \geq 2$, writing $J := j^*(l)$, if $J \geq l+1$ then $H(l, J-1) \pmod{9} \subseteq \{4^J, 4^{J+1}\} \pmod{9}$.*

Proof. As in Theorem 17's proof, every witness of $x \in D(l, W^*(l))$ at width $W^*(l) + 1$ has top exponent $\beta_0 = W^*(l) + 1$; modulo 9, only the rank-0 and rank-1 terms survive (rank $i \geq 2$ contributes a multiple of 3^2), so $x \equiv 2^{\beta_0} + 3 \cdot 2^{\beta_1} \pmod{9}$ for the witness's second-highest exponent β_1 . By (5), elements of $H(l, J - 1)$ are $y = 2^{J-1-l}x$, so $y \equiv 2^{J-1-l+\beta_0} + 3 \cdot 2^{J-1-l+\beta_1} \pmod{9}$; using $\beta_0 = W^*(l) + 1 = J + l - 1$, the first term is $2^{2J-2} = 4^{J-1} \pmod{9}$, fixed. The second term, $3 \cdot 2^m \pmod{9}$ for $m := J - 1 - l + \beta_1$, depends only on the parity of m (since $2^m \pmod{3}$ has period 2), taking exactly the two values 3 or 6 $\pmod{9}$ regardless of β_1 's actual value; a direct check gives $\{4^{J-1} + 3, 4^{J-1} + 6\} \equiv \{4^J, 4^{J+1}\} \pmod{9}$ for every J . So $y \pmod{9} \in \{4^J, 4^{J+1}\}$. $\square$

Empirical Result 21 (Exclusion of the lower class). The class actually occupied is 4^{J+1} , never 4^J , at every level checked, $l = 2, \dots, 16$. Excluding 4^J in general is the sharpest concrete open question this part of the investigation produced: Proposition 20 places no constraint on β_1 itself, only on the parity-dependent outcome, so ruling out one of the two parities needs a different argument.

Proposition 22 (Sufficiency of bounded maxrun). *If $\text{maxrun}(H(l, l + 1))$ is bounded by a constant C for all l , then $j^*(l) \leq l + C + 1$.*

Proof. Corollary 12 at $j = l + 1$. $\square$

The converse is not established, and the natural argument for it does not work: reading Corollary 12 downward gives only $\text{maxrun}(H(l, l + 1)) \geq j^*(l) - l - 1$, a *lower* bound on maxrun, which does not force maxrun to stay bounded when $j^*(l) - l$ does. Proving the converse would need an upper bound on maxrun in terms of $j^*(l) - l$, equivalent to a tightness statement for Corollary 12 that Section 8's corner-redundancy question addresses and that is false in the strong, per-witness form: recall 1547 from Section 6.

A naive expectation, that a random set of holdouts at the observed density would have doubling chains far longer than the flat ceiling of 3–4 measured at every level $l = 5, \dots, 22$, does not survive contact with the actual, rapidly declining density of $H(l, l + 1)$. An independence model matched to that actual density tracks the observed sequence within about one unit at every level, including both the one rise and the one fall in it: $\text{maxrun}(H(l, l + 1))$ is typical for the set's density under that model. This does not, by itself, say anything about whether $\text{maxrun}(H(l, l + 1))$ is bounded; Proposition 22 is the only established link between that question and $j^*(l)$'s rate.

Lemma 23 (One-step width identity). *Let $\text{Corner}(l, W + 1)$ be the set of values realized by a witness tuple $\beta_0 > \dots > \beta_{l-1}$ with $\beta_0 = W + 1$ and $\beta_{l-1} = 0$ simultaneously. Then*

$$U(l, W + 1) = U(l, W) \cup 2U(l, W) \cup \text{Corner}(l, W + 1).$$

Proof. $\supseteq$: $U(l, W) \subseteq U(l, W + 1)$ by monotonicity, $2U(l, W) \subseteq U(l, W + 1)$ by the shift argument of Proposition 16's proof, and $\text{Corner}(l, W + 1) \subseteq U(l, W + 1)$ since a corner witness is, in particular, a witness at width $W + 1$.

$\subseteq$: let $u \in U(l, W + 1)$, witnessed by $W + 1 \geq \beta_0 > \dots > \beta_{l-1} \geq 0$. If $\beta_0 \leq W$, the tuple already witnesses u at width W , so $u \in U(l, W)$. If $\beta_0 = W + 1$ and $\beta_{l-1} \geq 1$, subtracting 1 from every exponent gives a decreasing tuple in $[0, W]$ (top exponent W , bottom exponent ≥ 0) witnessing $u/2$, so $u \in 2U(l, W)$. If $\beta_0 = W + 1$ and $\beta_{l-1} = 0$, this is exactly the corner case, so $u \in \text{Corner}(l, W + 1)$. These three cases are exhaustive. $\square$

At every width W in the range $2l + 1 \leq W \leq j^*(l) + l - 1$, checked exhaustively for $l = 3, \dots, 13$ (the upper end is the largest width the exact table reaches at each level), $\text{Corner}(l, W + 1)$ contributes nothing new: every corner witness's value is already covered without it. Call this the *corner-redundancy* property at width W ; whether it holds at every width $W \geq 2l + 1$ in general, beyond the levels checked, is open. It also holds, checked directly, at the one smaller width $W = 2l$ itself for $l = 3, 4, 5, 6$; from $l = 7$ through 13 it is checked to fail there. The same exhaustive check also covers every width strictly below $2l$, down to $W = l - 1$, at every $l = 3, \dots, 13$, though no result

in this paper depends on that wider range. $W = 2l$ is exactly $j = l + 1$, Proposition 22's own case; corner-redundancy at that width is known, one way or the other, at every level $l = 3, \dots, 13$, holding only for $l = 3, 4, 5, 6$. What remains open is Proposition 22's own converse itself: corner-redundancy's failure at $W = 2l$ for $l \geq 7$ does not settle it either way, it only means this particular mechanism does not prove it there.

Proposition 24 (Corner-redundancy implies tightness). *Fix j with $l + 2 \leq j \leq j^*(l)$. If corner-redundancy holds at every width $2l + 1 \leq W \leq j^*(l) + l - 1$, then $j^*(l) = j + \text{maxrun}(H(l, j))$ exactly.*

Proof. Corner-redundancy at width W says $\text{Corner}(l, W + 1) \subseteq U(l, W) \cup 2U(l, W)$, so Lemma 23 gives $U(l, W + 1) = U(l, W) \cup 2U(l, W)$. Multiplication by 2 is a bijection of $(\mathbb{Z}/3^l\mathbb{Z})^*$, so it commutes with complementation and with intersection; taking complements turns this into $D(l, W + 1) = D(l, W) \cap 2D(l, W)$, and pushing through the shift in (5) (itself a bijection, for the same reason) turns the width-level equality at $W = j' + l - 1$ into an equality of holdout sets, $H(l, j' + 1) = 2H(l, j') \cap 4H(l, j')$, sharpening Theorem 10's inclusion to an identity whenever corner-redundancy holds at that width.

Claim: for every j' with $j \leq j' < j^*(l)$, $\text{maxrun}(H(l, j' + 1)) = \text{maxrun}(H(l, j')) - 1$, given corner-redundancy at width $j' + l - 1$. Since $j' < j^*(l)$, minimality of $j^*(l)$ gives $H(l, j') \neq \emptyset$, so $s := \text{maxrun}(H(l, j')) \geq 1$ throughout. The inequality $\leq$ is unconditional: if $H(l, j' + 1) = \emptyset$ it reads $0 \leq s - 1$, true since $s \geq 1$; otherwise a maximal chain of length $\text{maxrun}(H(l, j' + 1)) \geq 1$ in $H(l, j' + 1)$ extends, by Corollary 11, to a chain one longer in $H(l, j')$. For $\geq$, fix a maximal chain $x, 2x, \dots, 2^{s-1}x \in H(l, j')$. For each $k = 2, \dots, s$, both $2^{k-1}x$ and $2^{k-2}x$ lie in this chain, so the equality $H(l, j' + 1) = 2H(l, j') \cap 4H(l, j')$ places $2^k x \in H(l, j' + 1)$; this gives the $s - 1$ consecutive powers $4x, 8x, \dots, 2^s x \in H(l, j' + 1)$, so $\text{maxrun}(H(l, j' + 1)) \geq s - 1$, proving the claim.

The widths this uses, $j' + l - 1$ for $j \leq j' < j^*(l)$, all lie in $[j + l - 1, j^*(l) + l - 2] \subseteq [2l + 1, j^*(l) + l - 1]$ (using $j \geq l + 2$), inside the hypothesis, regardless of $\text{maxrun}(H(l, j))$'s actual value. Telescoping the claim over every step from $j' = j$ to $j' = j^*(l) - 1$ (a fixed number of steps, $j^*(l) - j$, fixed before any of this) gives $\text{maxrun}(H(l, j^*(l))) = \text{maxrun}(H(l, j)) - (j^*(l) - j)$. The left side is 0 ($H(l, j^*(l)) = \emptyset$ by definition), so $\text{maxrun}(H(l, j)) = j^*(l) - j$, i.e. $j^*(l) = j + \text{maxrun}(H(l, j))$. $\square$

Since corner-redundancy is verified at every width the proof above uses for $l = 3, \dots, 13$, Proposition 24 is unconditional at those levels: Empirical Result 13's equality is a proven theorem there, for every j with $l + 2 \leq j \leq j^*(l)$. In general, whether this implication's hypothesis holds is exactly the corner-redundancy question above.

At $l = 3, 4, 5, 6$, where corner-redundancy also holds at the boundary width $W = 2l$, the same argument (starting the telescoping at $j' = l + 1$ instead of $j' \geq l + 2$) proves the boundary case too: $j^*(l) = (l + 1) + \text{maxrun}(H(l, l + 1))$ exactly, verified directly ($l = 3$: $4 + 2 = 6$; $l = 4$: $5 + 2 = 7$; $l = 5$: $6 + 3 = 9$; $l = 6$: $7 + 3 = 10$). This is Proposition 22's own converse, established at these four levels specifically. At $l = 7, \dots, 13$, corner-redundancy at $W = 2l$ is checked to fail, so this mechanism does not reach the converse there; beyond $l = 13$, corner-redundancy at $W = 2l$ is itself unchecked. The converse remains open at every $l \geq 7$.

8. DISCUSSION

Two named questions summarize what stands between this paper's results and a sharper theorem. Excluding the lower class in Proposition 20's containment (Empirical Result 21) would turn the observed single-class pattern into a fully proven statement. Proving corner-redundancy at every width $W \geq 2l + 1$ would, by Proposition 24, make Corollary 12's bootstrap tight for every budget $l + 2 \leq j \leq j^*(l)$ in general, not merely at the levels $l = 3, \dots, 13$ where corner-redundancy is already verified; it would not, by itself, settle Proposition 22's converse at every level, which turns

on the boundary budget $j = l + 1$ ($W = 2l$), established above only at $l = 3, 4, 5, 6$, where corner-redundancy holds there too; at $l = 7, \dots, 13$, where it is checked to fail there, the converse remains open by this mechanism, and beyond $l = 13$ both the property and the converse are open. Neither question is, on its face, as hard as the Weak Covering Conjecture itself; whether either is much easier is not known.

A finite exponential sum built from strictly decreasing, exponentially weighted digit tuples is a natural object beyond this specific covering problem: finite-population L -statistics at non-CLT frequencies and weights, Littlewood–Offord theory for structured coefficient sequences, and the analytic side of the Collatz literature via Tao’s discussion are all adjacent areas. Findings here may recur there under different names. Generic cancellation, even at its best conceivable strength, cannot match a bound already conditionally established here by other means. Exceptional mass does not concentrate on a sparse set at the accessible scale. And an extremal arithmetic obstruction, pointed to by both holdout rarity and a phase-randomization null diagnostic, is exactly what magnitude-only methods miss. Tao’s later, unrelated result that almost every Collatz orbit attains an almost bounded value [3] works with a first-passage random variable over a different probabilistic model of the same map and does not bear on covering directly, but is the clearest instance of a new kind of technique changing what counts as tractable in this area, the kind of shift a resolution of either open question above would need. Predecessor-density bounds on the $3n + 1$ orbit structure itself, most developed in Krasikov and Lagarias’s difference-inequality approach [4] (a lower bound of $x^{0.84}$ on the count of integers below x whose orbit reaches 1), target the same broad question as Wirsching’s own reduction from covering to positive predecessor density, by a different mechanism: orbit-level difference inequalities rather than the covering structure of one finite family $R_{j-1,j}$. Nothing in that literature currently transfers to $K(l)$ or $j^*(l)$ directly, but the shared target makes a transfer more plausible than an unrelated result would, and whether one exists is itself an open question.

The Weak Covering Conjecture itself, and the exact asymptotic rate of $e(l)$, remain open.

9. CODE AND DATA AVAILABILITY

The repository <https://github.com/faculdade/weak-covering-conjecture> holds the code and data behind every computation reported here, one folder per section. It contains the covering search that produced Table 1, the model comparison behind Table 2, the mean-payoff-game solver, the exponential-sum computations of Section 5, and the holdout computations of Sections 6 and 7, together with the stored holdout sets at the final budgets of each level $l = 5, \dots, 21$; the $\text{maxrun}(H(l, l + 1))$ computation behind Proposition 22’s discussion, which reaches $l = 22$, lives as a directly runnable tool in its own folder. The twelve certificates (σ_k, ρ_k, h_k) behind Table 3 are stored there with a verifier that re-checks each of them against the inequality in Theorem 5’s proof, without re-running the solver. Each folder’s README gives what it checks, the command, the output to expect, and what the run costs. Most checks take seconds; the three largest levels of Table 1 do not.

Acknowledgments. AI assistance was used for textual review and translation.

REFERENCES

- [1] G. J. Wirsching, *The Dynamical System Generated by the $3n + 1$ Function*, Lecture Notes in Mathematics, vol. 1681, Springer, Berlin, 1998.
- [2] G. J. Wirsching, *On the problem of positive predecessor density in $3n + 1$ dynamics*, Discrete Contin. Dyn. Syst. **9** (2003), no. 3, 771–787.
- [3] T. Tao, *Almost all orbits of the Collatz map attain almost bounded values*, Forum Math. Pi **10** (2022), Paper No. e12, 1–56; arXiv:1909.03562.
- [4] I. Krasikov and J. C. Lagarias, *Bounds for the $3x + 1$ problem using difference inequalities*, Acta Arith. **109** (2003), no. 3, 237–258.

- [5] A. Ehrenfeucht and J. Mycielski, *Positional strategies for mean payoff games*, Internat. J. Game Theory **8** (1979), no. 2, 109–113.
- [6] T. Meyerovitch and A. Young, *Rationality and computability of the covering radius for sofic shifts*, arXiv:2603.21449 (2026).
- [7] T. Tao, *The Collatz conjecture, Littlewood–Offord theory, and powers of 2 and 3*, blog post, terry-tao.wordpress.com, August 25, 2011.