

# WIRSCHING’S WEAK COVERING CONJECTURE: AN EXTENDED COMPUTATION, A CONDITIONAL BOUND, AND THE STRUCTURE OF WHAT RESISTS IT

RENATO AUGUSTO TAVARES  
Universidade Federal de Goiás  
rat@discente.ufg.br

ORCID: <https://orcid.org/0009-0002-0196-3311>

**ABSTRACT.** Wirsching’s 1998 Weak Covering Conjecture asks for a sub-exponential  $K(l)$  such that a large enough set  $R_{j-1,j} \subset \mathbb{Z}_3^*$ , built from strictly decreasing exponent tuples, is forced to cover every unit residue modulo  $3^l$ . We extend the exact computation of  $j^*(l)$ , the smallest such covering budget, from  $l = 20$  to  $l = 23$ , and compare four growth models for  $e(l) := j^*(l) - l \log_4 3$  against the extended table: a constant model is disfavored at every statistic checked, but logarithmic, square-root, and slow-linear growth remain statistically indistinguishable from one another. We give the best bound known to us,  $j^*(l) \leq \frac{9}{8}l + O(1)$ , via a mean-payoff-game relaxation of the covering budget whose proof rests on one computationally verified, not generally proven, correspondence. We then examine the natural Fourier approach to a tighter bound: a uniform square-root-cancellation argument is capped below even this bound, a direct computation at one accessible level finds the bulk of the relevant exponential sum’s mass sitting outside any sparse exceptional set, and the residues that actually fail to be covered sit at extremal cells that a phase-randomization experiment shows are decided by phase structure that magnitude alone does not capture. We prove a theorem on the uncovered residues themselves: the set left uncovered at a budget is contained in both twice and four times the set left uncovered at the previous budget, giving an explicit recursive upper bound on  $j^*(l)$  that is empirically exact at every budget checked. We falsify the natural cost-1 form of a local repair rule this theorem might suggest, with explicit witness-level exhaustive counts, prove that every last-uncovered residue is  $\equiv 1 \pmod{3}$  whenever  $j^*(l) \geq l + 1$  (every computed level but one), and report several further exact identities, verified but not proven, governing the last residues to be covered. Two precisely stated open questions remain, each narrower than the conjecture itself. The Weak Covering Conjecture remains open.

**Keywords:** Collatz map, 3-adic dynamics, covering conjecture, mean-payoff games, exponential sums, Littlewood–Offord theory.

**2020 Mathematics Subject Classification:** 11B50, 11K70, 91A50, 11L07, 37A44.

## 1. INTRODUCTION

Wirsching’s 1998 monograph [1] studies the density of predecessor sets under the  $3n + 1$  map through an averaging construction on the 3-adic integers. A limiting Markov chain built from that construction has an invariant density, and Wirsching’s own route to positive predecessor density passes through a covering question about a family of finite subsets of  $\mathbb{Z}_3^*$ , the 3-adic units. For integers  $j, k \geq 0$ , write

$$(1) \quad R_{j,k} := \left\{ \sum_{i=0}^j 2^{a_i} 3^i \mid j+k \geq a_0 > a_1 > \cdots > a_j \geq 0 \right\} \subset \mathbb{Z}_3^*.$$

Each element of  $R_{j,k}$  is built from a strictly decreasing sequence of  $j + 1$  “digit positions” chosen from  $\{0, \dots, j + k\}$ : the  $i$ -th chosen position, counting from the top, contributes  $2^{a_i}$  scaled by  $3^i$ . Wirsching’s Weak Covering Conjecture (Conjecture 3.9, Chapter V, [1]) asks for a sub-exponential  $K(l)$  such that  $|R_{j-1,j}| \geq K(l) \cdot 3^l$  already forces  $R_{j-1,j}$  to cover every unit residue modulo  $3^l$ ; the cardinality on the left is that of the tuple family itself, before reduction, which by Proposition 1’s bijection equals  $\binom{2j}{j}$ , not the (a priori much smaller) size of its image modulo  $3^l$ . Writing  $j^*(l)$

for the smallest  $j$  such that  $R_{j-1,j}$  covers  $(\mathbb{Z}/3^l\mathbb{Z})^*$  exactly, the conjecture’s forcing form and the growth rate of  $j^*(l)$  coincide once coverage, once achieved, persists at every larger budget: this holds for  $j \geq l$  (Theorem 9, Section 6), which covers every level in Table 1, but is not established below  $l$ , the regime the conjectured rate would eventually enter. With that proviso, the conjecture concerns the growth rate of

$$(2) \quad e(l) := j^*(l) - l \log_4 3,$$

the excess of the true covering budget over the leading-order rate a cardinality argument alone predicts ( $|R_{j-1,j}| = \binom{2j}{j} \sim 4^j$ , so matching  $2 \cdot 3^{l-1}$  residues needs  $j \sim l \log_4 3$  at leading order). At the threshold,  $K(l) \asymp |R_{j^*(l)-1,j^*(l)}|/3^l \asymp 4^{j^*(l)}/3^l = 4^{e(l)}$ , so  $K(l)$  sub-exponential is equivalent to  $e(l) = o(l)$ : the conjecture is exactly the claim that  $e(l)$  grows slower than linearly, not just that it grows.

An earlier, unpublished manuscript of the author computed  $j^*(l)$  exactly for  $l = 1, \dots, 20$  by direct enumeration and reported the resulting table without a growth-rate analysis. This paper does three things. First, it extends that table to  $l = 23$  and gives a proper statistical comparison of candidate growth models for  $e(l)$  against the full range (Sections 2 and 3). Second, it gives the best upper bound on  $j^*(l)$  we have found, via a reformulation of the covering budget as a mean-payoff game (Section 4). Third, and at greatest length, it reports a sustained attempt to close the gap to the conjectured rate directly, first by the natural Fourier approach, and, once that approach is shown to fail for a precise and checkable reason, by a direct study of the residues that actually resist coverage (Sections 5 and 6). That second attempt does not settle the conjecture. It produces a theorem on those residues, a falsification of the cost-1 form of a repair rule the theorem might suggest, and two precisely named open questions, each narrower than the conjecture itself and worth a dedicated attack in its own right.

## 2. THE COMPUTATION

Table 1 gives  $j^*(l)$  and  $e(l)$  for  $l = 1, \dots, 23$ . The values for  $l \leq 20$  reproduce the earlier manuscript’s table exactly; they were independently recomputed here by a from-scratch Rust reimplement of the covering search (a bitset rotation dynamic program over residues modulo  $3^l$ ), validated against the original Python implementation at every  $l \leq 20$  before being trusted for the extension. The values for  $l = 21, 22, 23$  are new.

| $l$ | $j^*(l)$ | $e(l)$ | $l$ | $j^*(l)$ | $e(l)$ | $l$ | $j^*(l)$ | $e(l)$ |
|-----|----------|--------|-----|----------|--------|-----|----------|--------|
| 1   | 1        | 0.208  | 9   | 13       | 5.868  | 17  | 21       | 7.528  |
| 2   | 4        | 2.415  | 10  | 15       | 7.075  | 18  | 22       | 7.735  |
| 3   | 6        | 3.623  | 11  | 16       | 7.283  | 19  | 23       | 7.943  |
| 4   | 7        | 3.830  | 12  | 17       | 7.490  | 20  | 24       | 8.150  |
| 5   | 9        | 5.038  | 13  | 18       | 7.698  | 21  | 25       | 8.358  |
| 6   | 10       | 5.245  | 14  | 19       | 7.905  | 22  | 26       | 8.565  |
| 7   | 11       | 5.453  | 15  | 20       | 8.113  | 23  | 27       | 8.773  |
| 8   | 12       | 5.660  | 16  | 20       | 7.320  |     |          |        |

TABLE 1.  $j^*(l)$  and  $e(l) = j^*(l) - l \log_4 3$ ,  $l = 1, \dots, 23$ .

The single plateau in the table,  $j^*(15) = j^*(16) = 20$ , is not an error:  $l = 15 \rightarrow 16$  is the only increment among the 22 steps from  $l = 1$  to  $l = 23$  at which the covering budget does not increase, and it is exactly the event Section 3’s plateau-frequency test measures.

$l = 21$  fits in physical RAM once a memory-wasteful parallelization step in the original implementation was corrected.  $l = 22$  required 500 GiB of swap, added to the machine specifically for this computation, and took approximately 10,749 s ( $\sim 3$  hours) with swap I/O as the dominant

cost, against 748 s for  $l = 21$  without swap.  $l = 23$  required approximately 263 GiB of state and 375,615 s ( $\sim 104$  hours); the successful run had no interruption, so a checkpoint/resume mechanism added after an earlier, unrelated policy-triggered kill of an interrupted attempt at the same level was never exercised on it. Three attempts at  $l = 24$ , made over the course of this project on the same machine, each failed before completion (two to memory-policy kills, one to an apparent out-of-memory condition with no surviving kernel log confirming it directly);  $l = 24$  is not attempted further, and the table stops at  $l = 23$ .

### 3. THE GROWTH OF $e(l)$

We compare four candidate functional forms for  $e(l)$ 's growth against the tail  $l = 10, \dots, 23$  ( $n = 14$ ; the same range choice as the earlier manuscript, now with three more points), each fit by ordinary least squares:  $e(l) = c$  (constant),  $e(l) = a + b \ln l$  (logarithmic),  $e(l) = a + b\sqrt{l}$  (square-root), and  $e(l) = a + bl$  (slow-linear). Table 2 gives the fit statistics.

| model       | $k$ | $\Delta\text{AIC}$ | $\Delta\text{BIC}$ | LOOCV RMSE |
|-------------|-----|--------------------|--------------------|------------|
| constant    | 1   | 16.091             | 15.452             | 0.5200     |
| logarithmic | 2   | 1.512              | 1.512              | 0.2991     |
| square-root | 2   | 0.722              | 0.722              | 0.2909     |
| slow-linear | 2   | 0.000              | 0.000              | 0.2837     |

TABLE 2. Model comparison on  $e(l)$ ,  $l = 10, \dots, 23$ .  $\Delta\text{AIC}/\text{BIC}$  relative to the best-fitting model (slow-linear).

A constant model is disfavored sharply and consistently:  $\Delta\text{AIC} = 16.09$ ,  $\Delta\text{BIC} = 15.45$ , both growing at every  $l$  added since  $l = 21$ . The three growth models remain statistically indistinguishable from one another by the conventional  $\Delta\text{AIC} > 2$  rule of thumb, with slow-linear taking a narrow lead on both AIC and leave-one-out cross-validation; the three candidate regressors ( $\log l$ ,  $\sqrt{l}$ ,  $l$ ) are themselves highly correlated over this range (pairwise correlation  $\geq 0.993$ ), so this ranking says more about how thin the margin currently is than about which functional form is correct. The stakes are not merely descriptive: by the identity noted after (2), a genuinely linear  $e(l)$  would make  $K(l)$  exponential, falsifying the conjecture outright rather than leaving it merely unproven. Fourteen highly correlated points cannot distinguish that possibility from logarithmic or square-root growth, so this comparison does not settle the question either way; it only says which growth rates remain live.

A direct combinatorial test tells the same story from a different angle. Among the 22 increments  $j^*(l+1) - j^*(l)$  for  $l = 1, \dots, 22$ , exactly one is zero (the  $l = 15 \rightarrow 16$  plateau noted above). The earlier manuscript's own table,  $l \leq 20$ , showed 3.9 plateaus expected in 19 increments under a constant- $e(l)$  rounding model; taking  $q = 3.9/19 \approx 0.2053$  as that model's plateau probability and testing the observed count against  $\text{Bin}(22, q)$ , one-sided, gives  $p = 0.0426$ , crossing the conventional 0.05 threshold for the first time with the  $l = 23$  point included; restricted to the same tail the model comparison above uses ( $n = 13$ ), the same test gives  $p \approx 0.22$ . We report both figures:  $e(l)$  is an exact, deterministic sequence, so neither AIC/BIC nor this test carries the sampling interpretation of a classical inference over noisy data. Both are best read descriptively, as "how well does a constant- $e(l)$  source fit the observed increment pattern."

A constant  $e(l)$  is disfavored statistically by every measure above. It is also, unconditionally, impossible.

**Proposition 1** (Unboundedness of  $e(l)$ ).  $e(l) \geq \frac{1}{4} \log_2 l - O(1)$ .

*Proof.* Distinct exponent tuples give distinct sums: since the  $a_i$  are distinct,  $a_{j-1}$  is the unique minimal one, so  $V = \sum_i 2^{a_i} 3^i$  has 2-adic valuation exactly  $a_{j-1}$  (every other term is divisible by

a strictly higher power of 2), recovering  $a_{j-1}$  from  $V$  directly. The full term  $2^{a_{j-1}}3^{j-1}$  is then subtracted off, and the process iterates on the remaining  $j-1$  terms. So  $R_{j-1,j}$ , in bijection with  $j$ -subsets of the  $2j$ -element set  $\{0, \dots, 2j-1\}$ , has exact cardinality  $\binom{2j}{j}$ . Covering  $(\mathbb{Z}/3^l\mathbb{Z})^*$ , of size  $2 \cdot 3^{l-1}$ , requires the domain to be at least as large as the codomain:  $\binom{2j}{j} \geq 2 \cdot 3^{l-1}$  is necessary for  $j \geq j^*(l)$ . The explicit bound  $\binom{2n}{n} \leq 4^n/\sqrt{3n+1}$ , valid for every  $n \geq 1$ , turns this into  $4^j \geq 2 \cdot 3^{l-1}\sqrt{3j+1}$ , hence

$$j \geq l \log_4 3 - \log_4 3 + \log_4 2 + \frac{1}{2} \log_4(3j+1).$$

Dropping the last, non-negative term already gives  $j \geq l \log_4 3 - O(1)$ , so  $3j+1 = \Omega(l)$ , hence  $\log_4(3j+1) \geq \log_4 l - O(1) = \frac{1}{2} \log_2 l - O(1)$ . Feeding this back into the displayed inequality gives the claimed bound.  $\square$

#### 4. A MEAN-PAYOFF-GAME UPPER BOUND

**The full-precision game.** Fix  $l \geq 1$ . A *full-precision play* from a unit  $z_0$  modulo  $3^l$  is a sequence of  $l$  moves: at step  $i = 0, \dots, l-1$ , from a unit  $z_i$  modulo  $3^{l-i}$ , a move of cost  $d_i \geq 0$  is *legal* when  $2^{d_i}z_i \equiv 1 \pmod{3}$ , and produces  $z_{i+1} := T_{d_i}(z_i) := (2^{d_i+1}z_i - 2)/3$ , which is then required to be a unit modulo  $3^{l-i-1}$  (automatic at  $i = l-1$ , where  $3^{l-i-1} = 1$ ). Write  $J := \sum_{i=0}^{l-1} d_i$  for the total cost.

**Lemma 2** (Grounding). *For a legal full-precision play from  $z_0$  with costs  $d_0, \dots, d_{l-1}$  and total cost  $J$ , define  $a_{l-1} := 0$  and  $a_{i-1} := a_i + d_i + 1$  for  $i = l-1, \dots, 1$  (so  $a_0 = J - d_0 + l - 1$ ). Then*

$$2^{J+l-1}z_0 \equiv \sum_{i=0}^{l-1} 2^{a_i}3^i \pmod{3^l}.$$

The right side is, by construction, an element of  $U(l, a_0)$  (Section 7) with bottom exponent  $a_{l-1} = 0$ .

*Proof.* By induction on the play length  $n$  ( $n = 1, \dots, l$ ), stated for a play of length  $n$  starting at a unit  $z$  modulo  $3^n$  (matching the lemma's notation at  $n = l$ ,  $z = z_0$ ). Write  $J_n := \sum_{i=0}^{n-1} d_i$  and  $a_{n-1} := 0$ ,  $a_{i-1} := a_i + d_i + 1$  for  $i = n-1, \dots, 1$ , as above.

Base case  $n = 1$ : the single move has  $2^{d_0}z \equiv 1 \pmod{3}$ , which is exactly  $2^{a_0+d_0}z \equiv 2^{a_0}3^0 \pmod{3}$  since  $a_0 = 0$  when  $n = 1$ .

Inductive step: suppose the claim holds at length  $n-1$  for the sub-play with costs  $d_1, \dots, d_{n-1}$  starting at  $z' := T_{d_0}(z)$ , a unit modulo  $3^{n-1}$ ; writing  $a'_i := a_{i+1}$  for this sub-play,  $a'_{n-2} = a_{n-1} = 0$  matches its base case, and its recursion  $a'_{i-1} = a'_i + d_{i+1} + 1$  matches  $a_i = a_{i+1} + d_{i+1} + 1$  (the original recursion at index  $i+1$ ), so the induction hypothesis reads

$$2^{a_1+d_1}z' \equiv \sum_{i=0}^{n-2} 2^{a_{i+1}}3^i \pmod{3^{n-1}}.$$

By definition of  $T_{d_0}$ ,  $2^{d_0+1}z \equiv 3z' + 2 \pmod{3^n}$ . Multiplying by the unit  $2^{a_1+d_1}$  gives

$$2^{a_1+d_1+d_0+1}z \equiv 3 \cdot 2^{a_1+d_1}z' + 2^{a_1+d_1+1} \pmod{3^n}.$$

Since  $a_0 = a_1 + d_1 + 1$  (the recursion at  $i = 1$ ), the left side is  $2^{a_0+d_0}z$  and the last term is  $2^{a_0}$ . Multiplying the induction hypothesis by 3 turns it, modulo  $3^n$ , into  $3 \cdot 2^{a_1+d_1}z' \equiv \sum_{i=1}^{n-1} 2^{a_i}3^i$ , so

$$2^{a_0+d_0}z \equiv 2^{a_0} + \sum_{i=1}^{n-1} 2^{a_i}3^i = \sum_{i=0}^{n-1} 2^{a_i}3^i \pmod{3^n},$$

the claim at length  $n$ , using  $a_0 + d_0 = J_n + n - 1$  (telescoping the recursion). Verified directly against every legal play for  $l = 1, \dots, 4$  as a check, with no exception, before being trusted.  $\square$

*Empirical Result 3* (The game's worst case is  $j^*(l)$ ).

$$j^*(l) = \max_{z_0} \min\{J : z_0 \text{ has a legal full-precision play of total cost } J\},$$

the maximum over units  $z_0$  modulo  $3^l$ . Verified against the known exact table of  $j^*(l)$  for  $l = 1, \dots, 12$ , with no discrepancy. Lemma 2 does not by itself establish this: the twist  $2^{J+l-1}$  in the lemma depends on  $z_0$  through its own optimal  $J$ , so it is not a single,  $z_0$ -independent relabeling of the unit group, and we do not have a general proof of the displayed identity.

**The window- $k$  relaxation.** A move of cost  $d \geq 0$  at  $z$  is *legal* when  $2^d z \equiv 1 \pmod{3}$ , as above, and *safe* at window  $k$  when, for  $m := 3^k$ , every one of the three residues  $z, z + m, z + 2m$  modulo  $3m$  maps under  $T_d$  back into a unit modulo  $m$ . A minimizing player, seeing only  $z \bmod 3^k$ , chooses a legal, safe  $d$ ; an adversary then reveals the hidden next 3-adic digit  $\epsilon \in \{0, 1, 2\}$ , forcing the true state to  $T_d(z + 3^k \epsilon) \bmod 3^k$ . Both requirements matter: legality alone can send a unit to 0 (take  $z = 1, d = 0$ :  $T_0(1) = 0$ ), and the safety filter is exactly what keeps the walk inside  $\mathbb{Z}_3^*$  regardless of which digit the adversary supplies. Legality depends on  $d$  only through  $d \bmod 2$  (since 2 has order 2 modulo 3), and 2 is a primitive root modulo  $3^{k+1}$  for every  $k \geq 0$  (order  $2 \cdot 3^k$ , by an elementary lifting-the-exponent induction), so  $2^{d+1} \bmod 3^{k+1}$ , and hence  $T_d(z) \bmod 3^k$  and the safety check itself, depend on  $d$  only through  $d \bmod (2 \cdot 3^k)$ : successor states repeat with this period as  $d$  grows, and every representative beyond the smallest in its residue class has strictly larger cost for the same effect, so is never used by an optimal policy. The action set at each state is therefore, without loss, the subset of  $\{0, \dots, 2 \cdot 3^k - 1\}$  that is legal and safe there, which is finite. That this subset is nonempty at every unit state has been verified computationally at every  $k = 3, \dots, 14$  used below, not proven in general here.

This restricted-visibility game is, by construction, a two-player *mean-payoff game* on the  $2 \cdot 3^{k-1}$  unit states modulo  $3^k$ , with the now finite legal, safe moves as the minimizer's edges, the adversary's three choices of  $\epsilon$  as the maximizer's edges, and payoff the long-run average of the move costs  $d$ . The classical theory of such games, initiated by Ehrenfeucht and Mycielski [2], guarantees a value and an optimal positional policy in general; Theorem 4 below needs none of that generality, only the twelve specific policies actually exhibited and directly checked in Table 3.

A window- $k$  policy, run against the actual digit sequence, gives a legal full-precision play at any  $z$  modulo  $3^l$ : for  $l \geq k$  the policy sees the true  $z \bmod 3^k$  directly at every step; for  $l < k$ , fix any lift of  $z$  to a residue modulo  $3^k$  and run the policy on the lift, which is exactly the situation safety was built for, with the play's own remaining digits standing in for the hidden ones, so the chosen move stays legal against the true, lower-precision state at every step (this regime,  $l < k$ , is part of what the policy-against-the-covering-DP check at  $k = 4, 7$  against  $l$  up to 14 already exercised, with zero illegal or unsafe moves).

**Theorem 4** (Window relaxation bound, conditional on Empirical Result 3). *For each  $k = 3, \dots, 14$ , the certified policy underlying Table 3 (a legal, safe move at every unit state modulo  $3^k$ , average cost  $\rho_k$ ) gives an explicit rational  $C_k$  with  $j^*(l) \leq \rho_k \cdot l + C_k$  for all  $l$ .*

*Proof.* Fix  $k$  in range. The policy  $\sigma_k$ , the value  $\rho_k$ , and a potential  $h_k : (\mathbb{Z}/3^k\mathbb{Z})^* \rightarrow \mathbb{Q}$  are computed by a nested Howard strategy-improvement solver and checked, state by state and digit by digit over the finite action set established above, against

$$d_{\sigma_k}(s) + h_k(s'(\epsilon)) \leq \rho_k + h_k(s)$$

at every state  $s$  and every adversary digit  $\epsilon$ , where  $s'(\epsilon)$  is the successor under  $\sigma_k$ 's chosen move; the matching adversary lower bound extracted from the same run certifies  $\rho_k$  exactly. Summing the inequality along any  $l$ -step play  $s_0, \dots, s_l$  telescopes to

$$\sum_{i=0}^{l-1} d_{\sigma_k}(s_i) \leq l\rho_k + h_k(s_0) - h_k(s_l) \leq l\rho_k + (\max h_k - \min h_k) =: l\rho_k + C_k,$$

regardless of which digits the adversary actually supplies, since the potential inequality holds for every  $\epsilon$  at every step. So  $\sigma_k$ , run as above, gives every unit  $z$  modulo  $3^l$  a legal full-precision play of total cost at most  $l\rho_k + C_k$ ; the minimum cost over all legal plays from  $z$  is therefore also at most  $l\rho_k + C_k$ . Taking the worst-case  $z$  and applying Empirical Result 3 gives  $j^*(l) \leq \rho_k l + C_k$ .  $\square$

$\rho_k$  and  $C_k$  are computed exactly, as rationals, by a nested Howard strategy-improvement solver, self-certified at every  $k$  via a matching adversary lower bound extracted from the same run (upper bound equals lower bound, so the computed rational equals  $\rho_k$  exactly), and cross-validated by an independent value-iteration solver and, at  $k = 4$  and  $k = 7$ , by running the resulting policy against the actual finite covering DP for every unit residue modulo  $3^l$ ,  $l$  up to 14, with zero illegal or unsafe moves encountered. Table 3 gives the values computed for  $k = 3, \dots, 14$ .

| $k$      | 3 | 4              | 5              | 6             | 7             | 8                | 9              | 10              | 11             | 12             | 13                 | 14             |
|----------|---|----------------|----------------|---------------|---------------|------------------|----------------|-----------------|----------------|----------------|--------------------|----------------|
| $\rho_k$ | 2 | $\frac{5}{3}$  | $\frac{3}{2}$  | $\frac{3}{2}$ | $\frac{7}{5}$ | $\frac{25}{19}$  | $\frac{5}{4}$  | $\frac{11}{9}$  | $\frac{6}{5}$  | $\frac{7}{6}$  | $\frac{119}{104}$  | $\frac{9}{8}$  |
| $C_k$    | 5 | $\frac{19}{3}$ | $\frac{15}{2}$ | 9             | 10            | $\frac{207}{19}$ | $\frac{47}{4}$ | $\frac{115}{9}$ | $\frac{69}{5}$ | $\frac{44}{3}$ | $\frac{1621}{104}$ | $\frac{33}{2}$ |

TABLE 3. The window-relaxation value  $\rho_k$  and constant  $C_k$ ,  $k = 3, \dots, 14$ .

At  $k = 14$ ,  $\rho_{14} = 9/8 = 1.125$  and  $C_{14} = 33/2$ , giving:

**Corollary 5.**  $j^*(l) \leq \frac{9}{8}l + O(1)$ .

This is, to our knowledge, the best bound on  $j^*(l)$ 's growth rate currently available, conditional on Empirical Result 3.  $\rho_k$  is observed to be non-increasing over the computed range  $k = 3, \dots, 14$  (Table 3); we have not verified this in general, and Corollary 5 uses only the single value  $k = 14$ . The general technique, representing a covering-type combinatorial problem as a mean-payoff game and computing exact rational relaxation bounds this way, is not new to this application: beyond Ehrenfeucht and Mycielski's foundational theorem [2] itself, a close and recent methodological precedent proves rationality and computability of the covering radius of a primitive sofic shift by essentially the same construction, applied to a different object [3]. We cite both as precedent for the framing; what is new here is its specific application to  $j^*(l)$ .

*Remark 6.* Fitting  $\rho_k = L + A/k$  by ordinary least squares to  $k = 10, \dots, 14$  gives  $L \approx 0.876$ , still well above the conjectured true rate  $\log_4 3 = 0.7925\dots$ ; the fit is sensitive to the range used (all of  $k = 3, \dots, 14$  gives  $L \approx 0.904$ ) and is not evidence of any specific limit. Only the direction  $j^*(l) \leq \rho_k l + C_k$ , proven above, is established. The reverse inequality,  $\inf_k \rho_k \leq \limsup_l j^*(l)/l$ , which would be needed to conclude anything about where  $\rho_k$  is heading, is not established, and there is evidence in the course of this work that the window relaxation may carry an irreducible gap above the true rate.

## 5. THE EXPONENTIAL SUM, AND WHY IT RESISTS A DIRECT ATTACK

Corollary 5's bound,  $1.125l$ , is well above the conjectured rate  $\log_4(3)l \approx 0.7925l$ . The natural route to closing that gap is Fourier-analytic. By orthogonality on  $\mathbb{Z}/3^l\mathbb{Z}$ , the number of tuples landing in residue class  $z$  is

$$(3) \quad N(z) = \frac{1}{3^l} \sum_{t=0}^{3^l-1} S(t) e\left(-\frac{tz}{3^l}\right), \quad S(t) := \sum_{j+k \geq a_0 > \dots > a_j \geq 0} e\left(\frac{t}{3^l} \sum_{i=0}^j 2^{a_i} 3^i\right),$$

and, since  $S(0) = |R_{j,k}| =: T$ , the triangle inequality gives  $N(z) > 0$  for every  $z$  as soon as  $\sum_{t \neq 0} |S(t)| < T$ ; a uniform bound  $|S(t)| = o(T/3^l)$  at every nonzero  $t$  would suffice, at a rate controlled directly by the exponent in the bound. This approach does not work, for three separate and mutually reinforcing reasons.



### 5.1. A hard ceiling on uniform cancellation.

**Proposition 7.** *Any proof of covering via a uniform bound  $|S(t)| \leq C\sqrt{T}$  for a fixed constant  $C$  (square-root cancellation), applied via the triangle inequality across all  $3^l - 1$  nonzero frequencies, requires  $j \geq 2 \log_4(3)l - O(1) \approx 1.585l - O(1)$ .*

*Proof.* Write  $T := |R_{j-1,j}| = \binom{2j}{j}$ . Positivity of  $N(z)$  for every  $z$  via  $\sum_{t \neq 0} |S(t)| < T$  combined with the uniform estimate  $|S(t)| \leq C\sqrt{T}$  over the  $3^l - 1$  nonzero frequencies forces  $(3^l - 1)C\sqrt{T} < T$ , so  $T > ((3^l - 1)C)^2$ , i.e.  $4^{j/\sqrt{3j+1}} \geq T > ((3^l - 1)C)^2$  (using the explicit bound from Proposition 1's proof), which gives  $j \geq 2 \log_4(3)l - O(1)$ .  $\square$

Proposition 7's threshold is *worse* than Corollary 5's bound: any approach that only delivers generic square-root cancellation across every frequency cannot even match what is already known by other means, let alone improve on it. This rules out the naive version of the Fourier approach cleanly, and directs attention to whether a finer, frequency-dependent treatment (a major/minor-arc split, in the language of the circle method) can do better.

**5.2. A computation: the exceptional mass does not stay small.** Numerically,  $|S(t)|$  is far from uniform: the frequencies of largest magnitude sit close to dyadic rationals of small denominator,  $t/3^l \approx p/2^r$  for small  $r$ . The natural repair for Proposition 7's ceiling is to excise these major-arc resonances and control the remainder by an averaged bound; whether that remainder stays sparse enough at the scale that matters is a question this section answers computationally, for the family  $R_{j-1,j}$  at  $l = 18$ ,  $j = 16$  (the first-order cardinality threshold,  $T := \binom{32}{16}$ ), with a negative result. Write  $\|S\|_1 := \sum_{\gcd(t,3)=1} |S(t)|/T$  for the normalized  $L^1$  mass over frequencies coprime to 3; here  $\|S\|_1 = 5226.01$ . A fixed threshold makes the visible spikes look sparse, but of the 8,014 frequencies with  $|S(t)|/T$  exceeding 0.001, their combined contribution to  $\|S\|_1$  is under 132, so over 97% of the mass at this single computed level lives in an exponentially numerous population of near-average-magnitude coefficients, not in a sparse exceptional set. This is evidence against a sparse-exceptional-set repair at accessible  $l$ ; it says nothing about every  $l$ .

The same clustering near dyadic rationals connects the problem to a 2011 discussion of Terence Tao's, applying Littlewood–Offord and Riesz-product methods to an object built from the same  $2^a 3^b$ -type exponential sums, in the closely related setting of a hypothetical non-trivial Collatz cycle [4]; his construction is structurally close to a boundary slice of  $R_{j,k}$  with the bottom exponent pinned at 0. His own conclusion, in his own words, is that “once one uses the transcendence theory bound, ... Littlewood–Offord techniques are no longer of much use,” and that he knows “of no way to get a nontrivial separation property between powers of 2 and powers of 3 other than via transcendence theory methods.”

**5.3. A direct test of the residues that actually fail.** The most direct test examines the residues that actually fail to be covered, one budget step before covering succeeds, i.e.  $H(l, j^*(l) - 1)$  (Section 6). For  $l = 10, \dots, 15$  this set has size 2, 1, 3, 1, 3, 1 respectively, tiny against a global occupancy mean  $\lambda = |R_{j^*(l)-2, j^*(l)-1}|/(2 \cdot 3^{l-1})$  ranging 1.0 to 3.7 over the same levels; each holdout sits in a cell whose local intensity, at every conductor-depth resolution  $r \leq 10$  checked, assigns it probability between  $e^{-23}$  and  $e^{-100}$  under a non-homogeneous Poisson model fit to that intensity. A separate diagnostic at the higher-occupancy pair  $(l, m) = (14, 16)$  (mean hit count 188.5, well past the covering threshold) holds every  $|S(t)|$  fixed and randomizes the phases of conjugate-symmetric pairs: the actual primitive-frequency max/RMS ratio is 15.79, while eight independent phase scrambles (fixed seed) give ratios only in  $[5.11, 5.24]$ , collapsing to what a generic-phase model predicts once the phases are scrambled away. At the levels checked, both the raw holdout rarity and the phase-scramble gap point the same way: coverage failure is decided by phase structure that magnitude-only methods (uniform bounds, averages, low moments of  $S(t)$ ) do not see.

**Proposition 8.**  $|S(3^{l-1})|/\binom{2m}{m} \rightarrow 1/\sqrt{3}$  as  $m \rightarrow \infty$ , for the family  $R_{m-1,m}$ , for every fixed  $l \geq 1$ .

*Proof sketch.*  $\sum_i 2^{a_i} 3^i \equiv 2^{a_0} \pmod{3}$ , so  $S(3^{l-1})/\binom{2m}{m}$  depends only on the limiting distribution of the top exponent  $a_0$ 's parity as  $m \rightarrow \infty$ , which converges to the mixture  $\frac{1}{3}\omega + \frac{2}{3}\omega^2$  of the two nontrivial cube roots of unity  $\omega = e^{2\pi i/3}$  (the limiting gap  $2m-1-a_0$  of the maximum of a random  $m$ -subset of  $\{0, \dots, 2m-1\}$  is geometric with parameter  $1/2$ , giving  $\Pr[a_0 \text{ even}] \rightarrow 1/3$ ). This mixture equals  $-\frac{1}{2} - \frac{i\sqrt{3}}{6}$ , of modulus  $1/\sqrt{3}$ .  $\square$

Uniform decay across every nonzero frequency, the premise a naive statement of the Fourier approach might suggest, is false at this explicit family of frequencies. The regime is off-diagonal:  $l$  is fixed while  $m \rightarrow \infty$ , far past the covering-relevant scale  $m \sim l \log_4(3)$  where  $l$  and  $m$  grow together; the proposition rules out uniform decay in this one fixed- $l$  limit, not along the diagonal a covering argument would actually use.

## 6. A THEOREM ON THE RESIDUES THAT RESIST COVERAGE

Write  $H(l, j)$  for the set of units modulo  $3^l$  not covered by  $R_{j-1,j}$ , so that  $j \geq j^*(l)$  exactly when  $H(l, j) = \emptyset$ . Every result in this section needs  $j \geq l$ ; at every computed level  $j^*(l) \geq l$ , with equality only at  $l = 1$  (Table 1), but the conjectured rate  $\log_4 3 < 1$  would eventually put  $j^*(l)$  below  $l$ , outside this section's reach. Within this range, emptiness persists: Theorem 9 below gives  $H(l, j) = \emptyset \Rightarrow H(l, j+1) \subseteq 2\emptyset \cap 4\emptyset = \emptyset$ , so “ $H(l, j) = \emptyset$  exactly when  $j \geq j^*(l)$ ” is consistent for every  $j \geq l$ . We make no claim about  $j < l$ .

**Theorem 9** (Holdout doubling inclusion). *For  $j \geq l$ ,  $H(l, j+1) \subseteq 2H(l, j) \cap 4H(l, j)$ .*

*Proof.*  $R_{j-1,j}$  is the set of values  $r = \sum_{i=0}^{j-1} 2^{b_i} 3^i$  over  $j$ -subsets  $\{b_0 > \dots > b_{j-1}\} \subseteq \{0, \dots, 2j-1\}$ , and  $R_{j,j+1}$  the same with  $(j+1)$ -subsets of  $\{0, \dots, 2j+1\}$ . Fix  $r \in R_{j-1,j}$  with exponent set  $S$ , and  $s \in \{1, 2\}$ . Shifting every exponent up by  $s$  gives  $S+s \subseteq \{s, \dots, 2j-1+s\} \subseteq \{1, \dots, 2j+1\}$ , a  $j$ -subset whose induced value is  $2^s r$  (a uniform shift of exponents preserves relative rank, so each term's power of 3 is unchanged). Since  $S+s$  contains no 0, adjoining 0 as the new smallest exponent gives a  $(j+1)$ -subset of  $\{0, \dots, 2j+1\}$ : it becomes the new rank- $j$  term, contributing  $2^0 \cdot 3^j = 3^j$ . So this witness represents  $2^s r + 3^j$ , and  $3^j \equiv 0 \pmod{3^l}$  once  $j \geq l$ . Hence  $2R_{j-1,j} \cup 4R_{j-1,j} \subseteq R_{j,j+1} \pmod{3^l}$  for  $j \geq l$ . So if  $z \in H(l, j+1)$ , i.e.  $z$  is not represented at budget  $j+1$ , then  $z$  lies in neither  $2R_{j-1,j}$  nor  $4R_{j-1,j}$ : both  $z/2$  and  $z/4$  are unrepresented at budget  $j$ , i.e.  $z \in 2H(l, j) \cap 4H(l, j)$ .  $\square$

**Corollary 10** (Chain contraction). *If  $x, 2x, \dots, 2^{t-1}x \in H(l, j+1)$  (all reduced modulo  $3^l$ ,  $t \geq 1$ ), then  $x/4, x/2, x, \dots, 2^{t-2}x \in H(l, j)$ , a chain of length  $t+1$ .*

*Proof.* Apply Theorem 9 to each of the  $t$  elements:  $2^i x \in H(l, j+1)$  gives  $2^{i-1}x, 2^{i-2}x \in H(l, j)$  for  $i = 0, \dots, t-1$ . The union of these pairs over all  $i$  is exactly  $\{x/4, x/2, x, \dots, 2^{t-2}x\}$ , a doubling chain of length  $t+1$ .  $\square$

**Corollary 11** (Run-length bootstrap). *Writing  $\text{maxrun}(H)$  for the greatest  $t$  such that  $H$  contains  $t$  pairwise distinct elements of the form  $x, 2x, \dots, 2^{t-1}x$  (all reduced modulo  $3^l$ ), with  $\text{maxrun}(\emptyset) := 0$ ,  $j^*(l) \leq j + \text{maxrun}(H(l, j))$  for every  $j \geq l$ .*

*Proof.* Let  $r := \text{maxrun}(H(l, j))$  and suppose, for contradiction,  $H(l, j+r) \neq \emptyset$ ; pick  $z \in H(l, j+r)$ , a chain of length 1 ( $t = 1$ ) in the sense of Corollary 10. Applying Corollary 10 once at each of the  $r$  budgets  $j+r-1, j+r-2, \dots, j$  (all  $\geq l$ , since  $j \geq l$ ), a chain of length  $t$  at budget  $j'+1$  becomes a chain of length  $t+1$  at budget  $j'$  each time, so after  $r$  applications the length-1 chain at  $j+r$  becomes a chain of length  $r+1$  in  $H(l, j)$ , contradicting  $\text{maxrun}(H(l, j)) = r$ . So  $H(l, j+r) = \emptyset$ , i.e.  $j^*(l) \leq j+r = j + \text{maxrun}(H(l, j))$ .  $\square$



*Empirical Result 12.* Corollary 11 is empirically exact,  $j^*(l) = j + \max\text{run}(H(l, j))$ , at every  $(l, j)$  with  $l + 1 \leq j \leq j^*(l)$  we were able to compute,  $l = 5, \dots, 20$  at every such budget and  $l$  up to 22 at the single budget  $j = l + 1$ , including a case with  $|H(l, j)| > 3 \times 10^6$ . (Past  $j^*(l)$ ,  $H(l, j) = \emptyset$  and  $\max\text{run}(\emptyset) = 0$ , so the displayed equality no longer holds there; the claim is only about budgets up to and including  $j^*(l)$ .)

The converse direction, that a doubling chain of the observed maximal length is not merely sufficient but necessary for survival to that budget, is not proven, and is false as an unrestricted statement:  $1547 \in 2H(7, 8) \cap 4H(7, 8)$ , yet  $1547 \notin H(7, 9)$  (both of 1547's two budget-9 witnesses span the full exponent range, the one case Theorem 9's argument does not directly control). We return to a restricted, still-open form of this converse in Section 7.

A natural strengthening of Theorem 9, that every individual witness admits a bounded-cost local repair, is also false in its most optimistic form.

Call the *Durfee depth* of a witness  $\{a_0 > \dots > a_{j-1}\} \in R_{j-1, j}$  at budget  $j$  the count  $\#\{i : a_i \geq j\}$ , the number of chosen exponents at or above the budget itself.

*Empirical Result 13* (Failure of the cost-1 local repair rule). At the budget transition  $(l, j) = (6, 10) \rightarrow (7, 11)$ , 100 of the 1458 children (residues newly required for coverage at the larger budget) have minimum achievable Durfee depth exceeding their parent's minimum depth by more than one, up to three; at  $(7, 11) \rightarrow (8, 12)$ , 332 of 4374 children do, likewise up to three. A correspondingly exhaustive check, measuring witness repair cost as the minimum number of exponents that must change between a child's cheapest witness and its cheapest parent witness (matched by the natural containment of exponent sets), finds distances up to 7 at both transitions, against a distance of 1 for the overwhelming majority. This is evidence against a repair cost of exactly 1 at these two transitions; it does not establish that no bounded repair cost exists.

This falsifies, in its originally proposed cost-1 form, a natural conjecture this project had earlier entertained (that coverage propagates via a uniform, cost-1 local repair rule budget by budget); its valid surviving content is Theorem 9 and its corollaries above.

## 7. FURTHER STRUCTURE OF THE LAST HOLDOUTS

The residues in  $H(l, j^*(l) - 1)$ , the last to be covered, satisfy several further exact laws. For  $W \geq l - 1$ , write

$$U(l, W) := \left\{ \sum_{i=0}^{l-1} 2^{\beta_i} 3^i \bmod 3^l \mid W \geq \beta_0 > \dots > \beta_{l-1} \geq 0 \right\},$$

an  $l$ -term analogue of  $R_{j-1, j}$  built from a fixed term count and a growing exponent ceiling  $W$ .

**Lemma 14** (Width scaling identity). *For  $j \geq l \geq 1$ ,*

$$(4) \quad R_{j-1, j} \equiv 2^{j-l} U(l, j+l-1) \pmod{3^l}.$$

*Proof.* Since  $3^i \equiv 0 \pmod{3^l}$  for  $i \geq l$ , the residue of  $r = \sum_{i=0}^{j-1} 2^{b_i} 3^i \in R_{j-1, j}$  modulo  $3^l$  depends only on its top  $l$  exponents  $b_0 > \dots > b_{l-1}$ . A decreasing  $l$ -tuple  $b_0 > \dots > b_{l-1} \geq 0$  with  $b_0 \leq 2j - 1$  extends to a full admissible  $j$ -tuple in  $\{0, \dots, 2j - 1\}$  (appending  $j - l$  further, strictly smaller, distinct exponents) exactly when  $b_{l-1} \geq j - l$ : append  $b_{l-1} - 1, \dots, b_{l-1} - (j - l)$  when this holds, and no extension exists when it does not, since fewer than  $j - l$  nonnegative integers lie below  $b_{l-1}$ . Writing  $c_i := b_i - (j - l)$  turns  $b_{l-1} \geq j - l$  into  $c_{l-1} \geq 0$  and  $b_0 \leq 2j - 1$  into  $c_0 \leq 2j - 1 - (j - l) = j + l - 1$ , and  $\sum_i 2^{b_i} 3^i = 2^{j-l} \sum_i 2^{c_i} 3^i$ . So the top- $l$ -exponent tuples realizable by  $R_{j-1, j}$  are exactly  $2^{j-l}$  times the decreasing  $l$ -tuples counted by  $U(l, j+l-1)$ , giving the claimed identity modulo  $3^l$ .  $\square$

In particular  $j^*(l)$  is the smallest  $j \geq l$  with  $U(l, j + l - 1) = (\mathbb{Z}/3^l\mathbb{Z})^*$ , and writing  $W^*(l) := j^*(l) + l - 2$  for the width one step before that,

$$(5) \quad H(l, j) = 2^{j-l} K(l, j + l - 1), \quad K(l, W) := (\mathbb{Z}/3^l\mathbb{Z})^* \setminus U(l, W).$$

**Theorem 15** (Last-holdout parity). *If  $j^*(l) \geq l+1$ , every element of  $H(l, j^*(l)-1)$  is  $\equiv 1 \pmod{3}$ . This covers  $l = 2, \dots, 23$  (Table 1); it fails only at  $l = 1$  ( $j^*(1) = 1$ , outside (5)’s range), and would fail again for any  $l$  at which the conjectured rate  $\log_4 3 < 1$  eventually pushes  $j^*(l)$  below  $l + 1$ .*

*Proof.*  $U(l, \cdot)$  is monotone increasing in  $W$ , so  $K(l, \cdot)$  is monotone decreasing; every  $x \in K(l, W^*(l))$  is, by minimality of  $j^*(l)$ , representable at width  $W^*(l) + 1$  but not at  $W^*(l)$ . Any witness for  $x$  at width  $W^*(l) + 1$  using only exponents  $\leq W^*(l)$  would already witness  $x$  at that smaller width, so every such witness has top exponent  $\beta_0 = W^*(l) + 1$ . Its contribution to the value, at rank 0, is  $2^{\beta_0} \cdot 3^0 = 2^{W^*(l)+1}$ , so  $x \equiv 2^{W^*(l)+1} \equiv (-1)^{W^*(l)+1} \pmod{3}$ . By (5), which needs  $j^*(l) - 1 \geq l$ , elements of  $H(l, j^*(l) - 1)$  are  $2^{j^*(l)-1-l}x$  for such  $x$ , hence

$$2^{j^*(l)-1-l}x \equiv (-1)^{j^*(l)-1-l} \cdot (-1)^{W^*(l)+1} = (-1)^{2j^*(l)-2} = 1 \pmod{3},$$

using  $W^*(l) + 1 = j^*(l) + l - 1$  and that  $2j^*(l) - 2$  is always even.  $\square$

*Empirical Result 16* (Near-extinction bijection).  $H(l, j^*(l) - 1) = 2 \cdot \{x \in H(l, j^*(l) - 2) : x \equiv 2 \pmod{3}\}$ , exactly, verified at every  $l = 2, \dots, 9$  against the exact table with no exception.

The same width mechanism that proves Theorem 15 governs this bijection one step earlier, through the second-highest exponent  $\beta_1$  of the extremal witness; unlike the top-exponent argument above, pinning  $\beta_1$  down needs the one-step width recursion  $U(l, W + 1) = U(l, W) \cup 2U(l, W) \cup \text{Corner}(l, W + 1)$  stated below; we omit the full derivation and record the identity as verified.

An earlier stage of this investigation reported a cross-level “ $\times 4$ -only” pattern relating  $H(l + 1, j^*(l + 1) - 1)$  to  $H(l, j^*(l) - 1)$ . Writing  $\pi_l$  for reduction modulo  $3^l$ , the precise congruence this suggests,  $x \equiv 4\pi_l(x) \pmod{3^{l+1}}$ , is impossible whenever  $\pi_l(x)$  is a unit and  $l \geq 2$ : reducing both sides modulo  $3^l$  forces  $\pi_l(x) \equiv 4\pi_l(x) \pmod{3^l}$ , i.e.  $3\pi_l(x) \equiv 0 \pmod{3^l}$ . Checked exhaustively at  $l = 2, \dots, 6$ ,  $\pi_l(x)$  never lies in  $H(l, j^*(l) - 1)$  for  $x \in H(l + 1, j^*(l + 1) - 1)$  in the first place, so the pattern has no witnessing instance at these levels either. No cross-level relation between  $H(l + 1, \cdot)$  and  $H(l, \cdot)$  is established by this paper.

*Empirical Result 17* (Mod-9 two-class bound). Writing  $J := j^*(l)$ , at every  $l = 3, \dots, 9$  checked,  $H(l, J - 1) \pmod{9} \subseteq \{4^J, 4^{J+1}\} \pmod{9}$ , and the class actually occupied is  $4^{J+1}$ , never  $4^J$ , at every level checked. The two-class containment follows the same mechanism as Theorem 15 one digit further (the mod-9 class of a threshold witness is  $2^{\beta_0} + 3 \cdot 2^{\beta_1} \pmod{9}$ , with  $\beta_0$  pinned and  $\beta_1$ ’s parity giving the two options); excluding  $4^J$  specifically is the sharpest concrete open question this part of the investigation produced.

**Proposition 18** (Sufficiency of bounded maxrun). *If  $\text{maxrun}(H(l, l + 1))$  is bounded by a constant  $C$  for all  $l$ , then  $j^*(l) \leq l + C + 1$ .*

*Proof.* Corollary 11 at  $j = l + 1$ .  $\square$

The converse is not established, and the natural argument for it does not work: reading Corollary 11 downward gives only  $\text{maxrun}(H(l, l + 1)) \geq j^*(l) - l - 1$ , a lower bound on maxrun, which does not force maxrun to stay bounded when  $j^*(l) - l$  does. Proving the converse would need an upper bound on maxrun in terms of  $j^*(l) - l$ , equivalent to a tightness statement for Corollary 11 that Section 8’s corner-redundancy question addresses and that is false in the strong, per-witness form: recall 1547 from Section 6.

A naive expectation, that a random set of holdouts at the observed density would have doubling chains far longer than the flat ceiling of 3–4 measured at every level  $l = 5, \dots, 22$ , does not survive contact with the actual, rapidly declining density of  $H(l, l + 1)$ . An independence model matched

to that actual density tracks the observed sequence within about one unit at every level, including both the one rise and the one fall in it:  $\text{maxrun}(H(l, l+1))$  is typical for the set's density under that model. This does not, by itself, say anything about whether  $\text{maxrun}(H(l, l+1))$  is bounded; Proposition 18 is the only established link between that question and  $j^*(l)$ 's rate.

The one-step width identity is exact:

$$U(l, W+1) = U(l, W) \cup 2U(l, W) \cup \text{Corner}(l, W+1),$$

where  $\text{Corner}(l, W+1)$  consists of the witnesses using both the top exponent  $W+1$  and the bottom exponent 0 simultaneously. At every width  $W \geq 2l+1$  checked exhaustively,  $l = 3, \dots, 13$ ,  $\text{Corner}(l, W+1)$  contributes nothing new: every corner witness's value is already covered without it. If this *corner-redundancy* property holds at every width  $W \geq 2l+1$  in general, it would make Corollary 11's bootstrap tight for every budget  $j \geq l+2$  (the width  $W = j+l-1$  ranges over  $W \geq 2l+1$  exactly there); it remains unproven, with the only known failures at the boundary width  $W = 2l$ , which is exactly  $j = l+1$ , Proposition 18's own case, so corner-redundancy at  $W \geq 2l+1$  does not by itself bear on that proposition's converse.

## 8. DISCUSSION

Two named questions summarize what stands between this paper's results and a sharper theorem. Deciding which of the two classes in Empirical Result 17 is excluded would turn the observed single-class pattern into a proven statement. Proving corner-redundancy at every width  $W \geq 2l+1$  would make Corollary 11's bootstrap tight for every budget  $j \geq l+2$ , not merely empirically exact at every computed instance; it would not, by itself, settle Proposition 18's converse, which turns on the single boundary budget  $j = l+1$  ( $W = 2l$ ), exactly where corner-redundancy is known to fail. Neither question is, on its face, as hard as the Weak Covering Conjecture itself; whether either is substantially easier is not known.

A finite exponential sum built from strictly decreasing, exponentially weighted digit tuples is a natural object beyond this specific covering problem: finite-population  $L$ -statistics at non-CLT frequencies and weights, Littlewood–Offord theory for structured coefficient sequences, and the analytic side of the Collatz literature via Tao's discussion are all adjacent areas. Three things found here may recur there under different names: generic cancellation capping out below what combinatorial counting already proves, exceptional mass that does not concentrate on a sparse set at the accessible scale, and coverage failure decided by an extremal arithmetic obstruction that a phase-randomization experiment showed magnitude-only methods miss.

The Weak Covering Conjecture itself, and the exact asymptotic rate of  $e(l)$ , remain open.

## 9. CODE AND DATA AVAILABILITY

The repository <https://github.com/faculdade/weak-covering-conjecture> holds the code and data behind every computation reported here, one folder per section. It contains the covering search that produced Table 1, the model comparison behind Table 2, the mean-payoff-game solver, the exponential-sum computations of Section 5, and the holdout computations of Sections 6 and 7, together with the stored holdout sets at the final budgets of each level  $l = 5, \dots, 21$ . The twelve certificates  $(\sigma_k, \rho_k, h_k)$  behind Table 3 are stored there with a verifier that re-checks each of them against the inequality in Theorem 4's proof, without re-running the solver. Each folder's README gives what it checks, the command, the output to expect, and what the run costs. Most checks take seconds; the three largest levels of Table 1 do not.

**Acknowledgments.** AI assistance was used for textual review and translation.

## REFERENCES

- [1] G. J. Wirsching, *The Dynamical System Generated by the  $3n + 1$  Function*, Lecture Notes in Mathematics, vol. 1681, Springer, Berlin, 1998.
- [2] A. Ehrenfeucht and J. Mycielski, *Positional strategies for mean payoff games*, Internat. J. Game Theory **8** (1979), no. 2, 109–113.
- [3] T. Meyerovitch and A. Young, *Rationality and computability of the covering radius for sofic shifts*, arXiv:2603.21449 (2026).
- [4] T. Tao, *The Collatz conjecture, Littlewood–Offord theory, and powers of 2 and 3*, blog post, terry-tao.wordpress.com, August 25, 2011.